

CMPE110 Lecture 21

Input / Output (I/O) 3

Heiner Litz

<https://canvas.ucsc.edu/courses/12652>



Announcements


Review



Bus



- A shared communication link that connects multiple devices
 - Single set of wires connects multiple “subsystems”
 - As opposed to a point to point link which only connects two components
- Wires connect in parallel
 - E.g., 32 bit bus has 32 wires of data
 - And a few additional wires for control






Increasing Bus Bandwidth

■ Block transfers


- Transfer multiple words in back-to-back bus cycles
- Only one address needs to be sent at the start
- Bus is not released until the last word is transferred
- Costs: Increased complexity and response time for pending requests





Increasing Bus Bandwidth

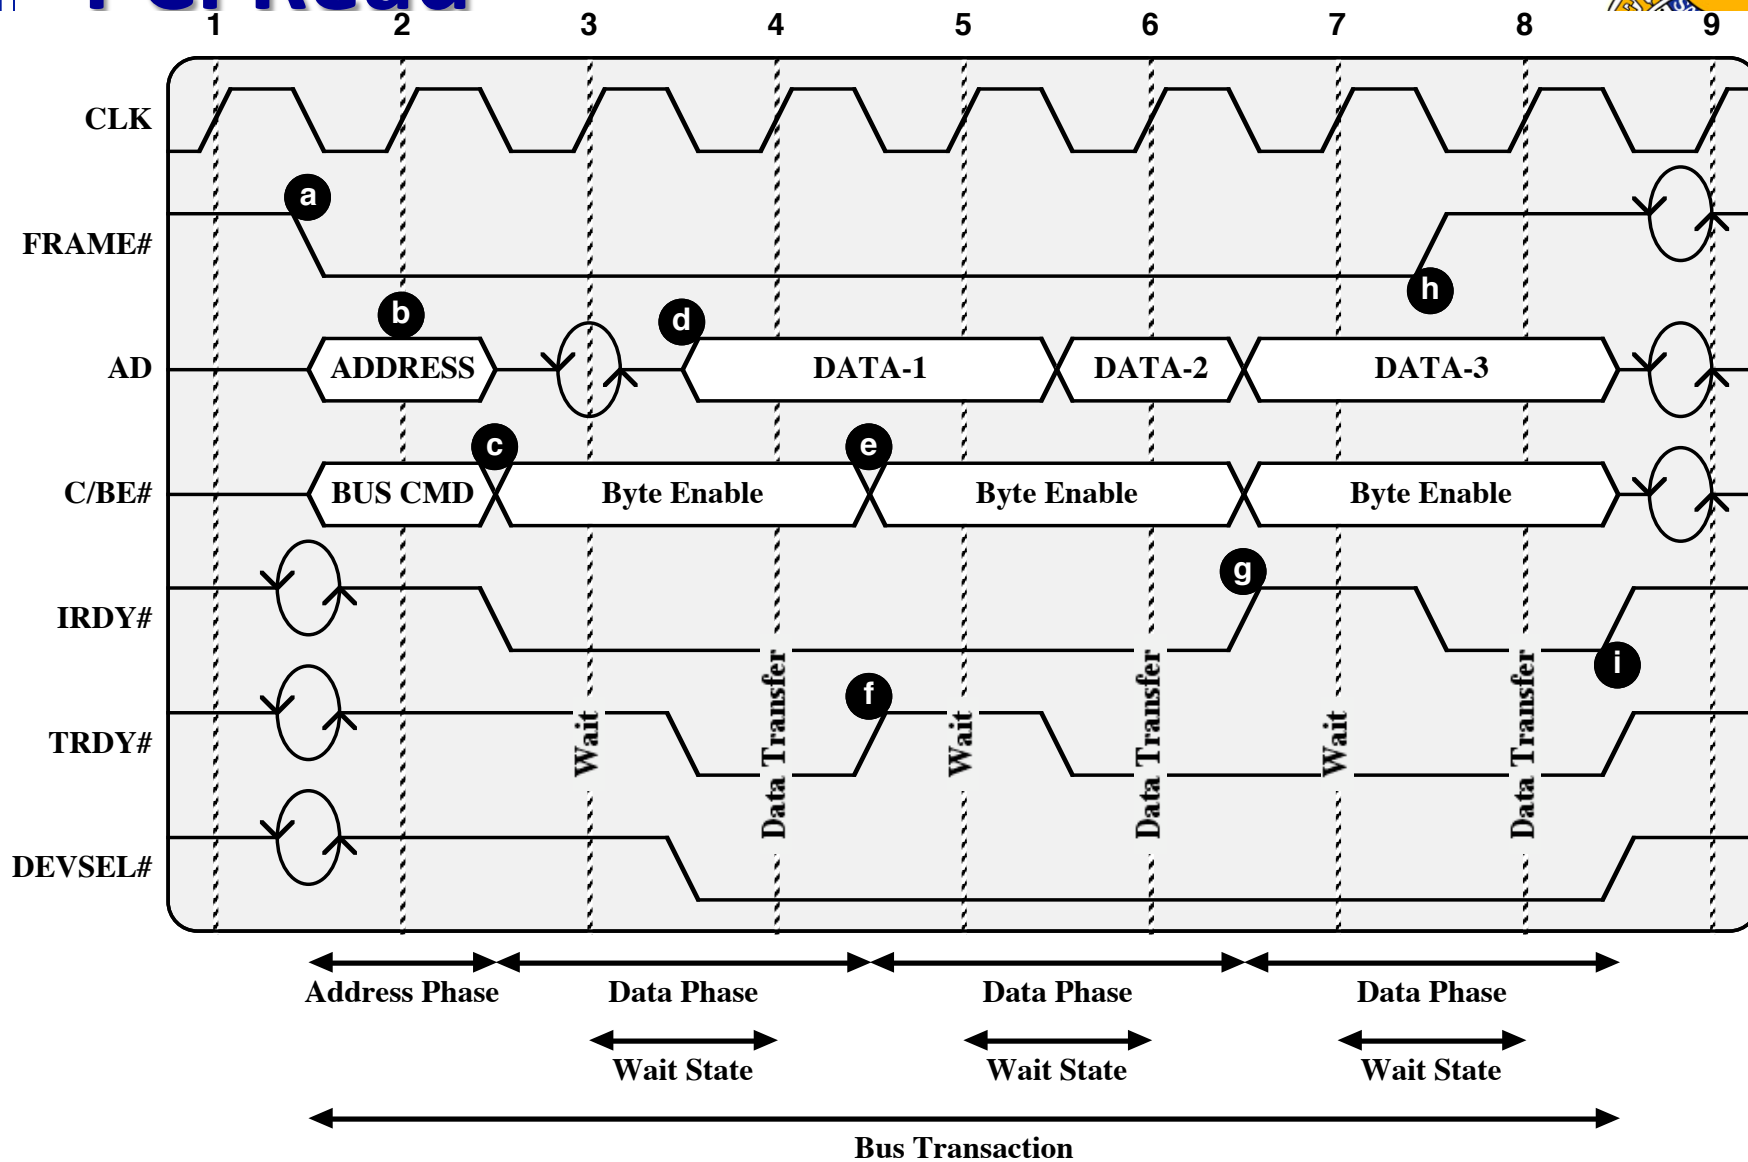
- Split transaction “pipelining the bus”
 - Free the bus during time between request and data transfer
 - Costs: Increased complexity and higher potential latency



Split transaction bus with separate Address and Data wires



PCI Read






Buses Revisited

- Advantages:
 - Simple
 - Broadcast
 - Serialization
 - Where is this useful? Tip: think multi-core
- Drawbacks: many ways of saying it's slow
 - Only one transaction at a time
 - Electrically ugly
 - Stubs and discontinuities limit operating speed
 - Global closed-loop handshakes




Alternative: Point-to-Point Links & Switches



Serial point-to-point advantages

- Faster links
- Fewer chip package pins
- Higher performance
- Switch keeps arbitration on chip



- Many bus standards are moving to serial, point to point
- 3GIO, PCI-Express(PCI)
- Serial ATA (IDE hard disk)
- AMD Hypertransport, Intel QPI versus Intel Front Side Bus (FSB)



PCI vs. PCI Express

- Same bus protocol
 - Same driver software
- PCI
 - 32–64 shared wires
 - Frequency: 33MHz – 133 MHz
 - Bandwidth: 132 MB/s – 1 GB/s
- PCI Express
 - 1-16 wires per direction
 - Frequency: 2.5GT/s (PCIe 2.0: 5GT/s, PCIe 3.0: 8GT/s, PCIe 4.0: 16GT/s)
 - Bandwidth: 250 MB/s per direction per wire (500MB/984MB/1669MB)
- PCI Express advantage
 - ~100x pin bandwidth
 - Multiple links (lanes) for more bandwidth

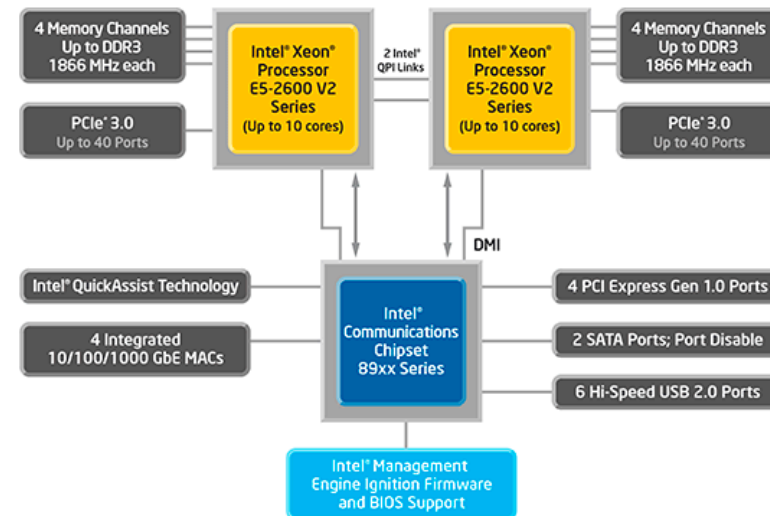


Buses Vs Switched Interconnects

- Both useful, both used
- Busses used mostly on-chip
 - Mostly when connecting very small number of components
 - But also off-chip sometimes, e.g. DDR-x for main memory
- Switch-based interconnects used on and off-chip
 - Particularly when connecting many components
 - But topology can also be simple (e.g., rings)



Hierarchies of Interconnects



- Modern systems have hierarchies of buses/interconnects
 - Each chip tracks address maps
 - Memory-mapped I/O requests forwarded down the hierarchy until they reach the proper device
 - Potentially with translation along the way



I/O Notification

- How does the OS or our program know that something interesting happened in the I/O device?
 - E.g., a packet arrived on the network interface
 - E.g., an error occurred while reading the disk




I/O Notification with Polling

- The OS periodically checks the status register
 - I/O device places information in a status register
- Advantages
 - Simple to implement
 - Processor is in control and does the work
- Disadvantage
 - Polling overhead and data transfer consume CPU time
 - Difficult to determine how frequently to poll



I/O with Polling

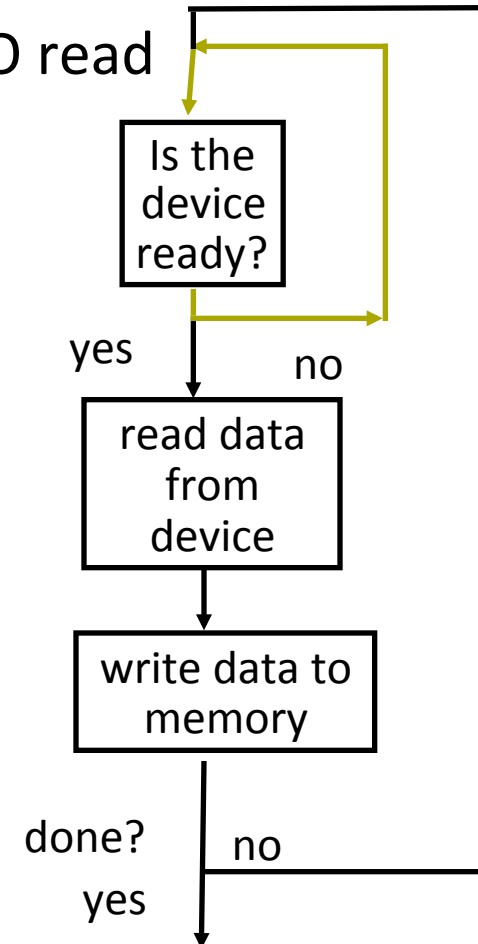
I/O write



Busy wait loop
not an efficient
way to use the CPU
unless the device
is very fast!

Processor may be
inefficient way to
transfer data

I/O read





I/O Notification with Interrupts


- When I/O device needs attention, it *interrupts* the processor
 - Just like an exception but due to an external source and asynchronous
- Cause register helps OS find the I/O device
 - Or vector interrupts in the case of x86
- Interrupts may have different priorities
 - Ex.: Network = high priority, keyboard = low priority
- Advantage
 - Execution is only halted during actual transfer
- Disadvantage
 - Software overhead of interrupt processing



I/O Interrupts

Interrupts Exceptions

Exception Vector Table



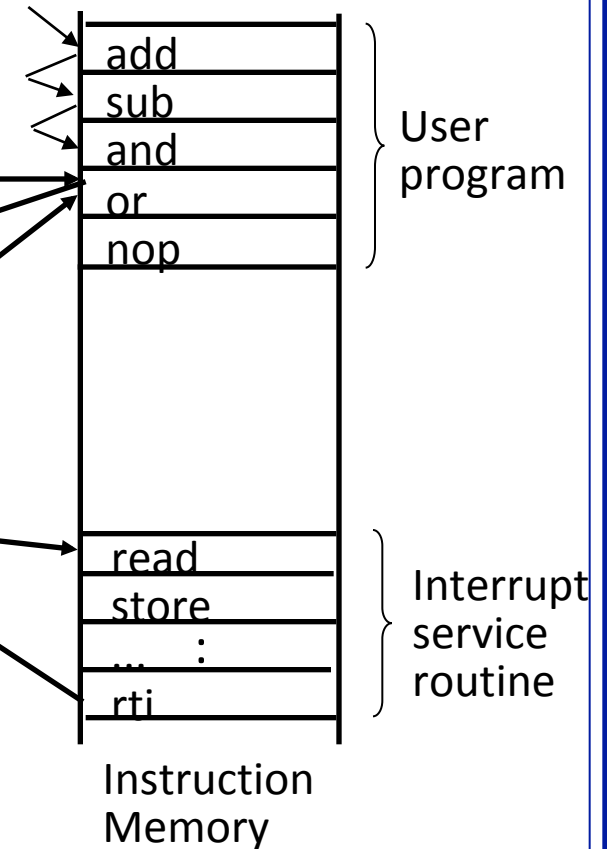
(contains PCs of ISRs)

(1) I/O interrupt

(2) Save state

(3) Interrupt
service addr

(4) Restore
state and return





Polling Vs. Interrupts

- Polling works great when
 - A device has a fixed I/O rate (e.g., keyboard or mouse)
 - The rate of interrupts is too high
 - The latency of interrupts is unacceptable
- Interrupts work best when
 - When devices can initiate their own I/O events
 - When I/O rate or latency is unpredictable
 - E.g. disk access
 - The overhead (CPU cycles) of polling is unacceptable



Large Data Transfers

- Simple approach: use ld/st instructions
 - SW moves *all* data between memory and I/O addresses
 - Simple and flexible, but wastes processor time
- Need a solution to allow data transfer to happen without the processor's involvement



Direct Memory Access (DMA)

- A custom engine for data movement
 - Transfer blocks of data to or from memory without CPU intervention
 - DMA engines are available in processor chips and I/O chips
 - DMA engine acts as a bus master in bus-based systems
- DMA engine interface
 - The DMA engine is an I/O device itself
 - Registers or memory to describe transfers
 - “From”, “to”, and “length” registers
 - Can have multiple of those or a queue
 - Registers for status and control

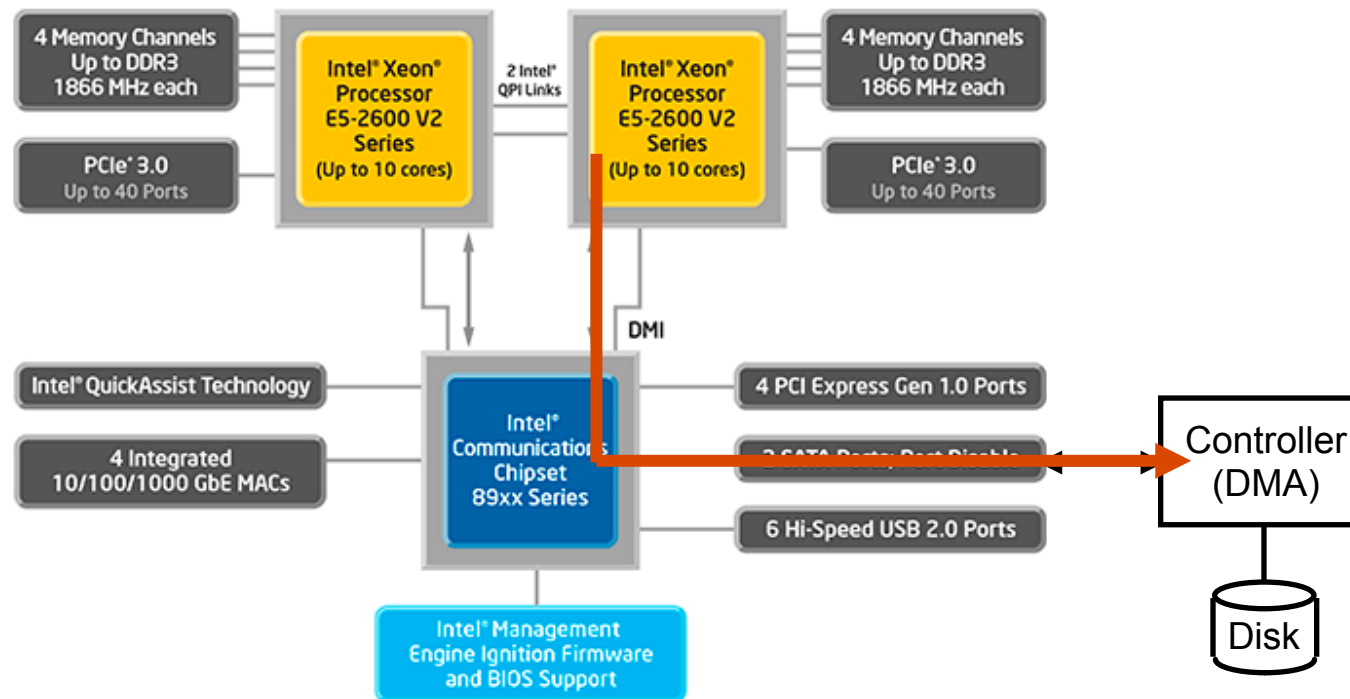


DMA Steps

- Processor sets up DMA by supplying
 - Identity of the device and the operation (read/write)
 - The memory address for source/destination
 - The number of bytes to transfer
- DMA engine starts by arbitrating for bus and then starting transfer when data is ready
- Notify processor when transfer is complete or on error
 - Typically with interrupt



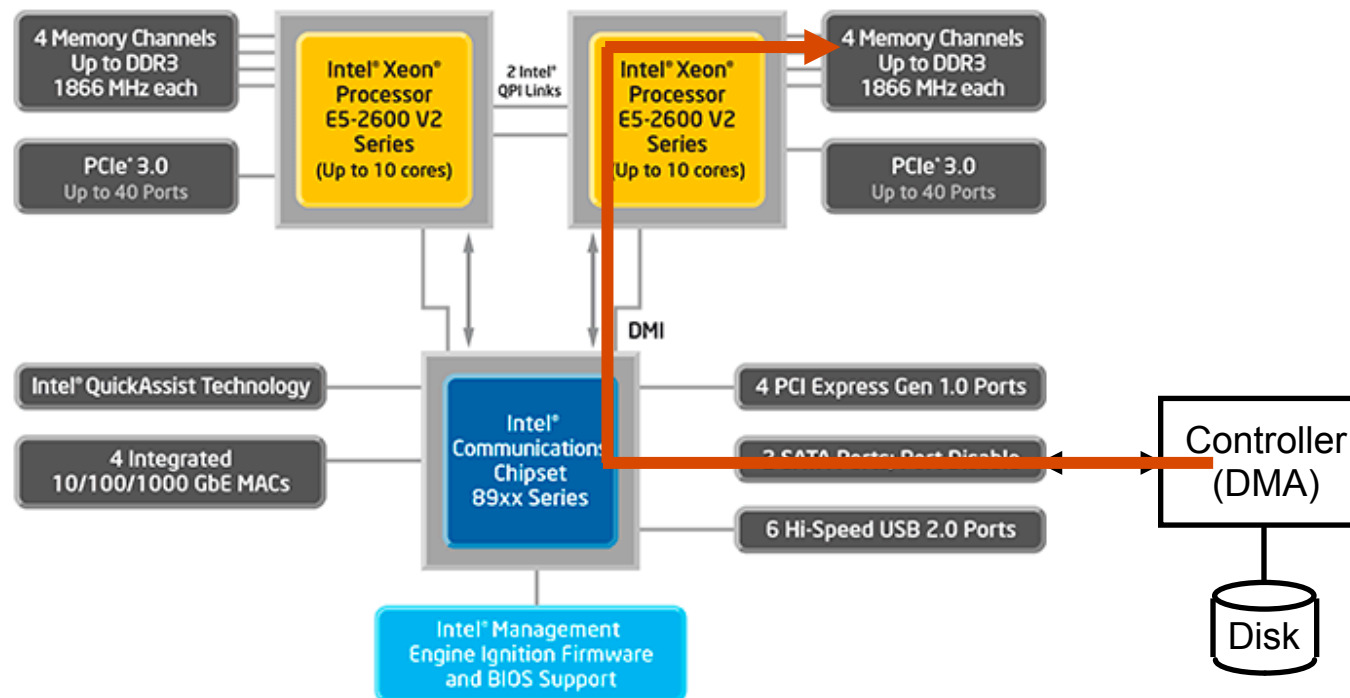
Reading a Disk Sector (1)



- CPU initiates a disk read by writing a command, logical block number, and a destination memory address to disk controller control registers.



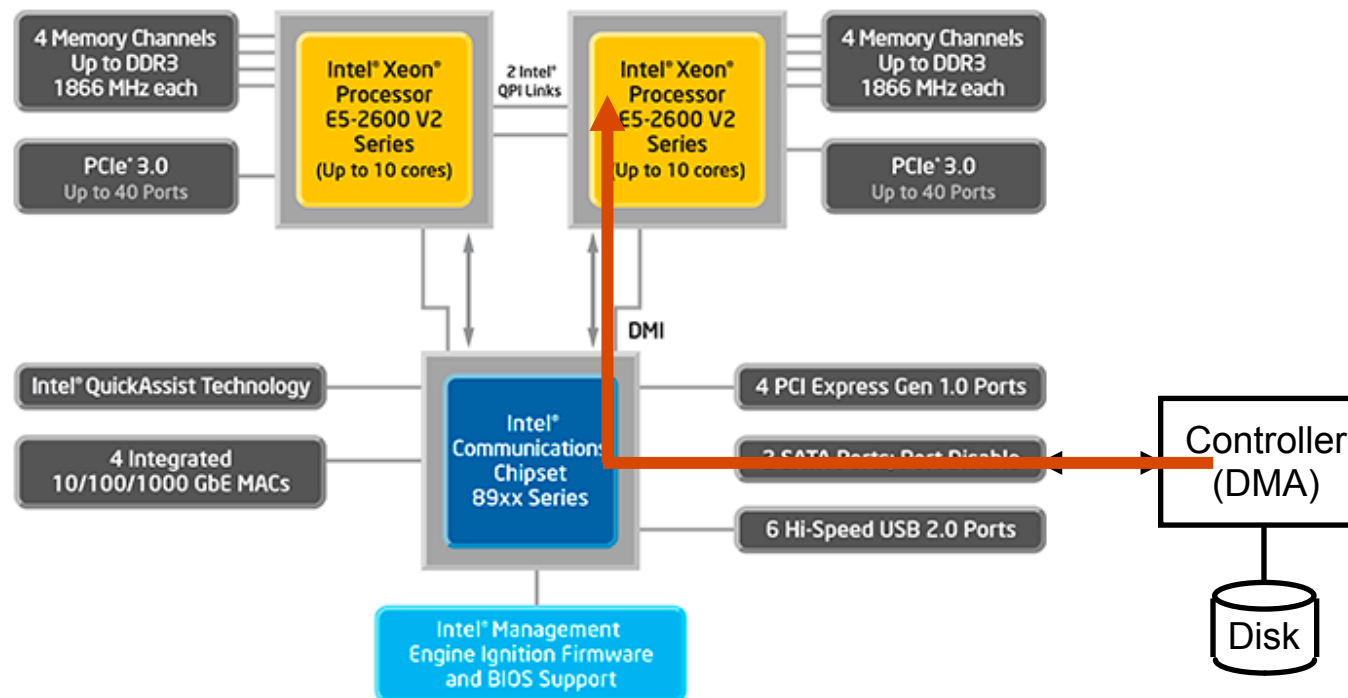
Reading a Disk Sector (2)



- Disk controller reads the sector and performs a direct memory access (DMA) transfer into main memory (Disk is Master).



Reading a Disk Sector (3)



- When the DMA transfer completes, the disk controller notifies the CPU with an *interrupt*



DMA Issues (1): Memory Pinning

- For DMA with physical addresses, memory pages must be pinned in DRAM - Why?
- OS should not page to disks pages involved with pending I/O



DMA Issues (2): Memory Translation

- If DMA uses physical addresses
 - Memory access across physical page boundaries may not correspond to contiguous virtual pages (or even the same application!)
- Solution 1: ≤ 1 page per DMA transfer
- Solution 1+: chain series of 1-page requests provided by the OS
 - Single interrupt at the end of the last DMA request in the chain
- Solution 2: DMA engine uses virtual addresses
 - Multi-page DMA requests are now easy
 - A TLB is necessary for the DMA engine
 - Who manages the TLB?

DMA Issues (3): Cache Coherence



- The data involved in DMA may reside in processor cache
 - If updated by I/O: must update or invalidate “old” cached copy
 - If read by I/O: must read latest value, which may be in the cache
 - Only a problem with write-back caches

- Another version of the “cache coherence” problem
 - Same problem in multi-core systems



DMA & Coherence

- Solution 1: OS flushes the cache before I/O reads or forces write-backs before I/O writes
 - Flush/write-back may involve selective addresses or whole cache
 - Can be done in software or with hardware (ISA) support
- Solution 2: Route memory accesses for I/O through the cache
 - Search the cache for copies and invalidate or write-back as needed
 - This hardware solution may impact performance negatively
 - While searching cache for I/O requests, it is not available to processor
 - Multi-level, inclusive caches make this easier
 - Processor searches L1 cache mostly (until it misses)
 - I/O requests search L2 cache mostly (until it finds a copy of interest)