

CMPE110 Lecture 25

Security and GPUs

Heiner Litz

<https://canvas.ucsc.edu/courses/12652>



Announcements

- Next Friday Review Session 2
 - Post your questions on Piazza

- Final Exam
 - Tuesday, June 12th
 - 12 – 3pm



Security



OOO Processors

Modern processors optimize for ILP

Characteristics

- Fetch multiple instructions per cycle (2-4, superscalar)

- Execute instructions out-of-order

 - Execution order dictated by true dependencies

 - Memory/registers updated in order due to precise exception

- Speculative

 - Branch prediction

 - Also speculate no exceptions (page faults, etc)

 - Since memory/registers updated in order this is ok

- Deep cache hierarchies (2 to 4 layers)

OOO Execution – Dynamic Scheduling



Execute instructions out-of-order

Fetch multiple instructions per cycle using branch prediction

Figure out which are independent and execute them in parallel

Example

```
add    $t0, $t1, $t2
or     $t3, $t0, $t2
sub    $t0, $t1, $t2
and    $t5, $t0, $t2
```

Superscalar + Dynamic scheduling

```
add    $t0, $t1, $t2      sub $t0, $t1, $t2
or     $t3, $t0, $t2      and $t5, $t0, $t2
```

Requires register renaming to make it work

Recent Security Attacks to OOO Processors



Based on slides from Mark Hill

Full set at <https://goo.gl/QcwqeP>

Key issues

- Our ISAs specify functionality but not timing

- Our processors have a lot of internal state

 - Known as microarchitecture state

- Attacks exploit timing behavior and internal state to bypass security checks

Side-Channel Attack: **PRIME** Secret in Micro-Arch



1. Prime micro-architectural state

- a. Repeatedly access array **train[]** to train branch predictor to expect access $< \text{bound}$
- b. Access all of array **save[]** to put it completely in a cache of size **SIZE**

Side-Channel Attack: **SAVE** Secret in Micro-Arch



2. Coerce processor into **speculatively executing** instructions that will be nullified to (a) find a secret & (b) save it in micro-architecture

```
branch (R1 >= bound) goto error ; Speculate not taken even if R1 >= bound
load R2 ← memory[train+R1]      ; Speculate to find SECRET outside of train[]
and R3 ← R2 && 0xffff            ; Speculate to convert SECRET bits into index
load R4 ← memory[save+SIZE+R3] ; Speculate to save SECRET by victimizing
memory[save+R3] since it aliases in cache with new access memory[save+SIZE+R3]
```

3. HW detects **mis-speculation**

Undoes architectural changes

Leaves cache (**micro-architecture**) changes (correct by Architecture 1.0)

Side-Channel Attack: **RECALL** Secret from Micro-Arch



Spy vs. Spy, Mad Magazine, 1960

4: Probe **time** to access each element of **save []** --
micro-architectural property;

If accessing **save [foo]** slow due to cache miss,
then SECRET is **foo**. A leak!

5: Repeat many times to obtain secret information
at some bandwidth. (More shifting/masking needed
to get all SECRET bits victimizing 64B cache lines)

Well-known as covert timing channel (1983)

Meltdown v. Spectre



MELTDOWN

Architecture

Intel, Apple

Entry

Must have code execution on the system

Method

Intel Privilege Escalation + Speculative Execution

Impact

Read kernel memory from user space

Action

Software patching



SPECTRE

Intel, Apple, ARM, AMD

Must have code execution on the system

Branch prediction + Speculative Execution

Read contents of memory from other users' running programs

Software patching (more nuanced)

Daniel Miessler 2018

Miessler Blog (<https://danielmiessler.com/blog/simple-explanation-difference-meltdown-spectre/>)

Meltdown



Can leak the contents of kernel memory at up to 500KB/s



```
1 ; rcx = kernel address
2 ; rbx = probe array
3 retry:
4 mov al, byte [rcx]
5 shl rax, 0xc
6 jz retry
7 mov rbx, qword [rbx + rax]
```

TRAP!! (not branch)
Under mis-speculation

Listing 2: The core instruction sequence of Meltdown. An inaccessible kernel address is moved to a register, raising an exception. The subsequent instructions are already executed out of order before the exception is raised, leaking the content of the kernel address through the indirect memory access.

Meltdown & Software



Bad: Meltdown operates with bug-free OS software

Good: Major commercial OSs patched for Meltdown ~January 2018

Idea: Don't map (much) of protected kernel address space in user process

Spectre



Classic side-channel attack w/ deep micro-arch info

- Many ways to prime prime micro-architecture
 - Branch predictor
 - Branch target buffer
- Multiple covert channels
 - Timing (cache, register file, functional unit contention)
 - Power consumption
 - EMI

Spectre Applicability



Exploit branch mis-prediction to let Javascript steal from Chrome browser

- Demonstrated Intel Haswell/Skylake, AMD Ryzen, & several ARM cores
- Many other existing designs vulnerable

Spectre Mitigation



Branch prediction

- SW: Suppress branch prediction “when important” with **mfence**, etc.
- Insert mfence manually into all software?
- What if we miss one?
- HW could auto-magically suppress branch prediction when appropriate (???)

Spectre Mitigation



Branch Target Buffer

- SW: Not clear. Disable hyper-threading, etc.?
- HW: Make micro-architecture state private to thread (not core or processor)

More generally: Hard to mitigate threats NOT YET DEFINED. We don't have a good answer to this yet!

GPUs





History of GPUs

- Early video cards
 - Frame buffer memory with address generation for video output
- 3D graphics processing
 - Originally high-end computers (e.g., SGI)
 - Moore's Law \Rightarrow lower cost, higher density
 - 3D graphics cards for PCs and game consoles
- Graphics Processing Units
 - Processors oriented to 3D graphics tasks
 - Vertex/pixel processing, shading, texture mapping, rasterization

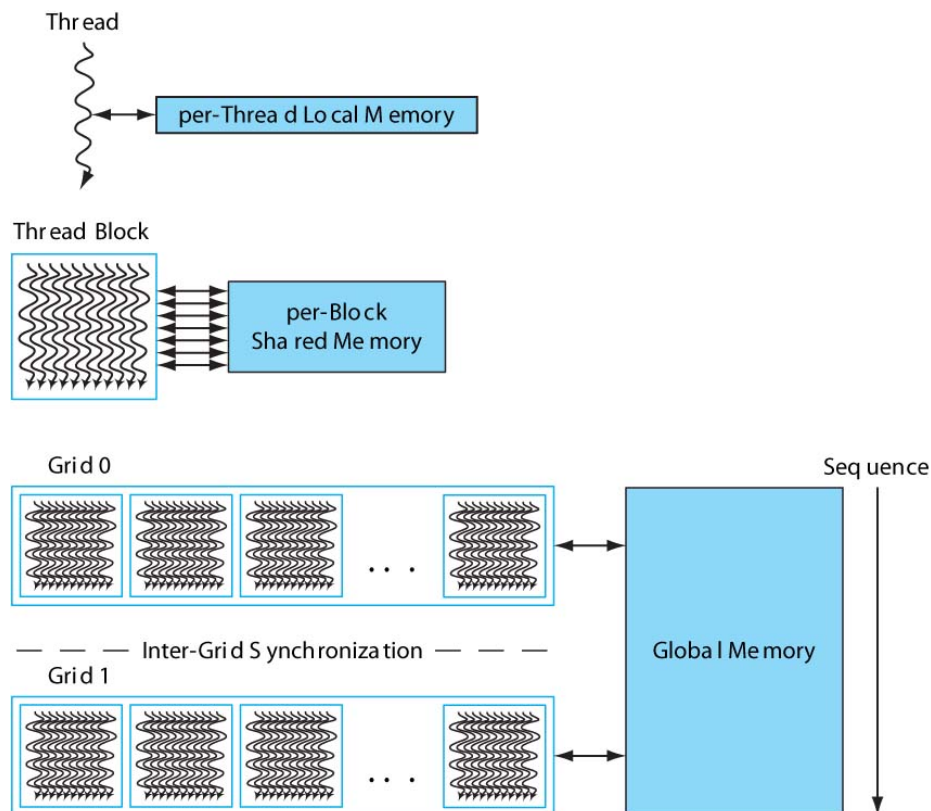


GPU Architectures

- Processing is highly data-parallel and highly multi-threaded
 - Use thread switching to hide memory latency
 - Less reliance on multi-level caches
 - Graphics memory is wide and high-bandwidth
- Trend toward general purpose GPUs
 - Heterogeneous CPU/GPU systems
 - CPU for sequential code, GPU for parallel code
- Programming languages/APIs
 - DirectX, OpenGL
 - C for Graphics (Cg), High Level Shader Language (HLSL)
 - Compute Unified Device Architecture (CUDA)



GPU Thread Model Software View



Single instruction multiple threads
(SIMT)

Parallel threads packed in blocks

- Also called Warps
- All threads execute in lockstep
- All threads perform same instruction
- Access to per-block shared memory
- Control flow divergence within a block with predication
- Switch thread block each cycle within a grid
- Can synchronize with barrier

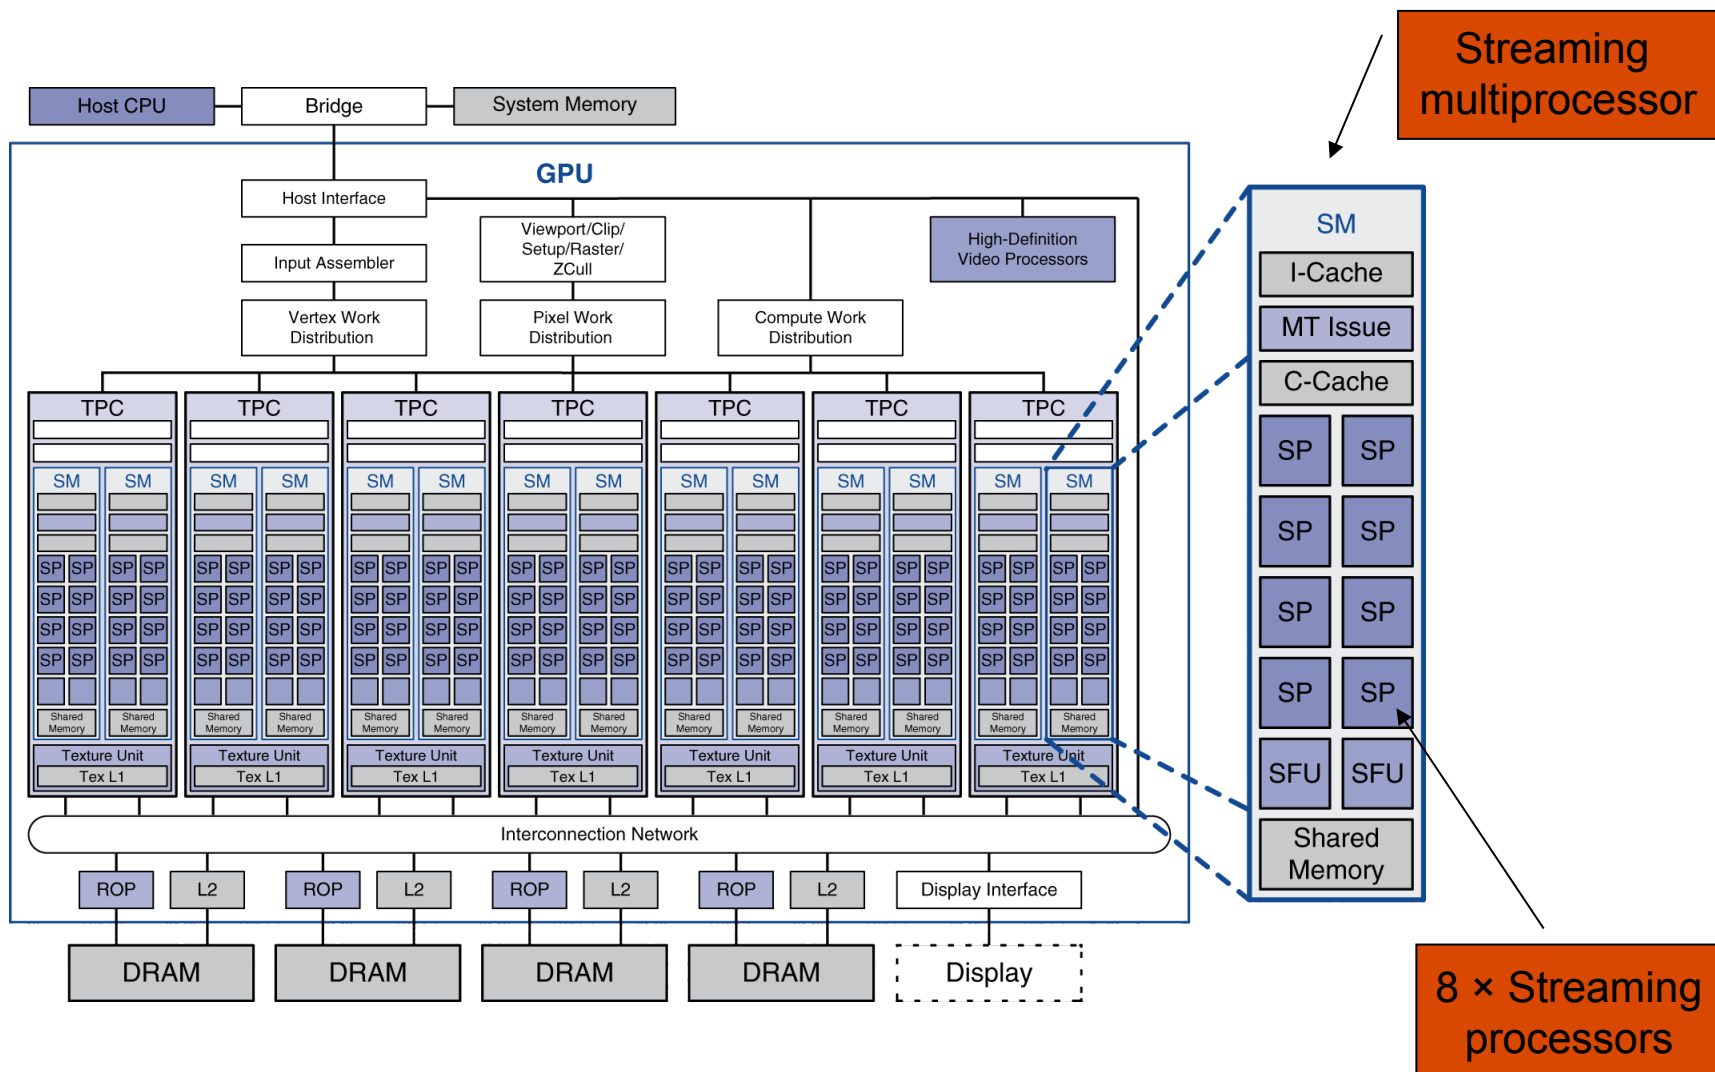


GPU Memory Hierarchy

- Registers
 - Needed to support large number of threads without spilling to main memory: 64K (vs 32 of RISC-V)
- Shared Memory
 - Statically shared among threads
 - Explicitly managed by programmer
- Hardware managed cache
 - Only in recent GPGPUs, mainly for power reduction
- DRAM
 - High Bandwidth, small, soldered on PCB



Example: NVIDIA Tesla



Putting GPUs into Perspective



Feature	Multicore with SIMD	GPU
SIMD processors	4 to 8	8 to 16
SIMD lanes/processor	2 to 4	8 to 16
Multithreading hardware support for SIMD threads	2 to 4	16 to 32
Typical ratio of single precision to double-precision performance	2:1	2:1
Largest cache size	8 MB	0.75 MB
Size of memory address	64-bit	64-bit
Size of main memory	8 GB to 256 GB	4 GB to 6 GB
Memory protection at level of page	Yes	Yes
Demand paging	Yes	No
Integrated scalar processor/SIMD processor	Yes	No
Cache coherent	Yes	No