

CMPE110 Lecture 11

Pipelining III

Heiner Litz

<https://canvas.ucsc.edu/courses/12652>




Announcements

Review





Forwarding Paths: Partial





Double Data Hazard

- Consider the sequence:

add **\$1**, \$1, \$2 Mem Fwd
sub **\$1**, **\$1**, \$3 Ex Fwd
or \$1, **\$1**, \$4

- Both hazards occur
 - Want to use the most recent result from the sub
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true



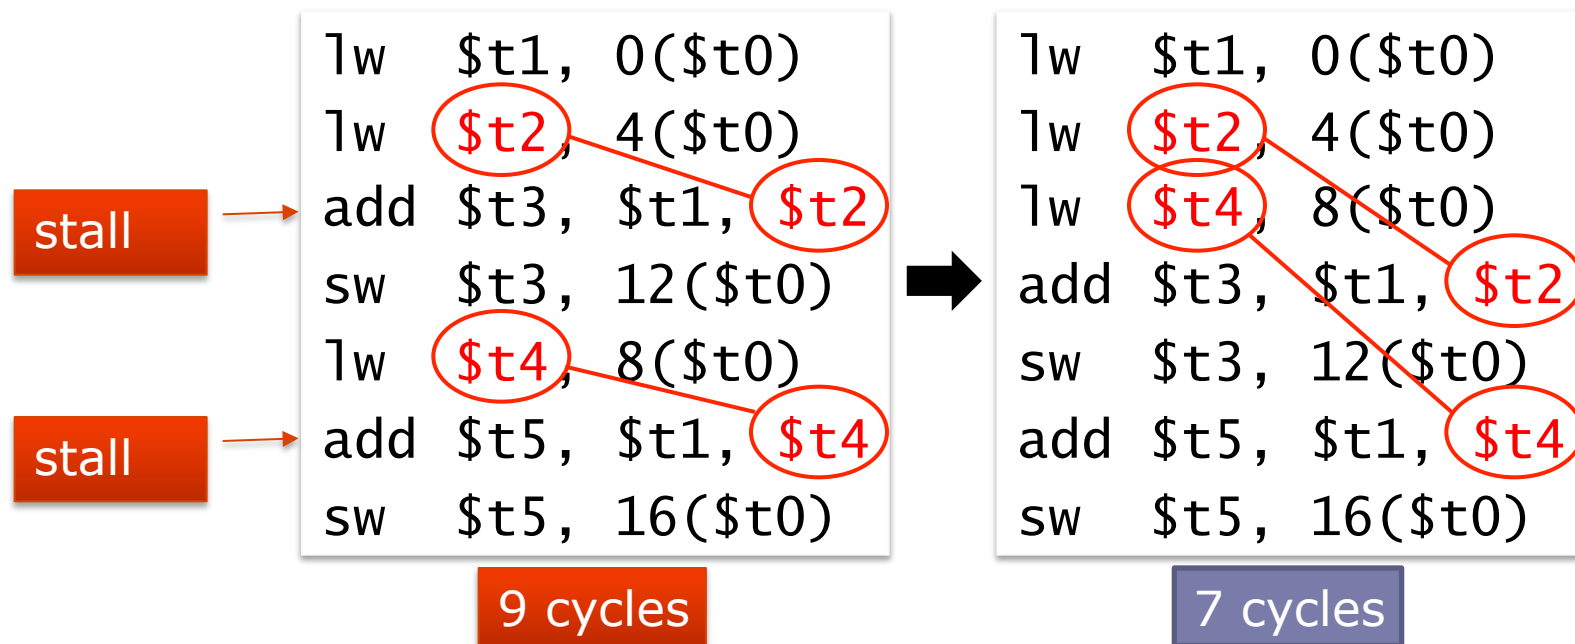
Forwarding Control (Revised)

- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd \neq 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd \neq 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs))
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd \neq 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd \neq 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRt))
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))
ForwardB = 01



Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for $A = B + E$; $C = B + F$;





Control Hazards




Reducing Branch Delay

- Minimize “bubble” slots
 - Move branch computation earlier in the pipeline
 - branch outcome: add comparator to ID stage
 - branch target: add adder to ID stage

- Predict branch not taken
 - if correct, no bubbles inserted
 - if wrong, flush pipe, inserting one bubble



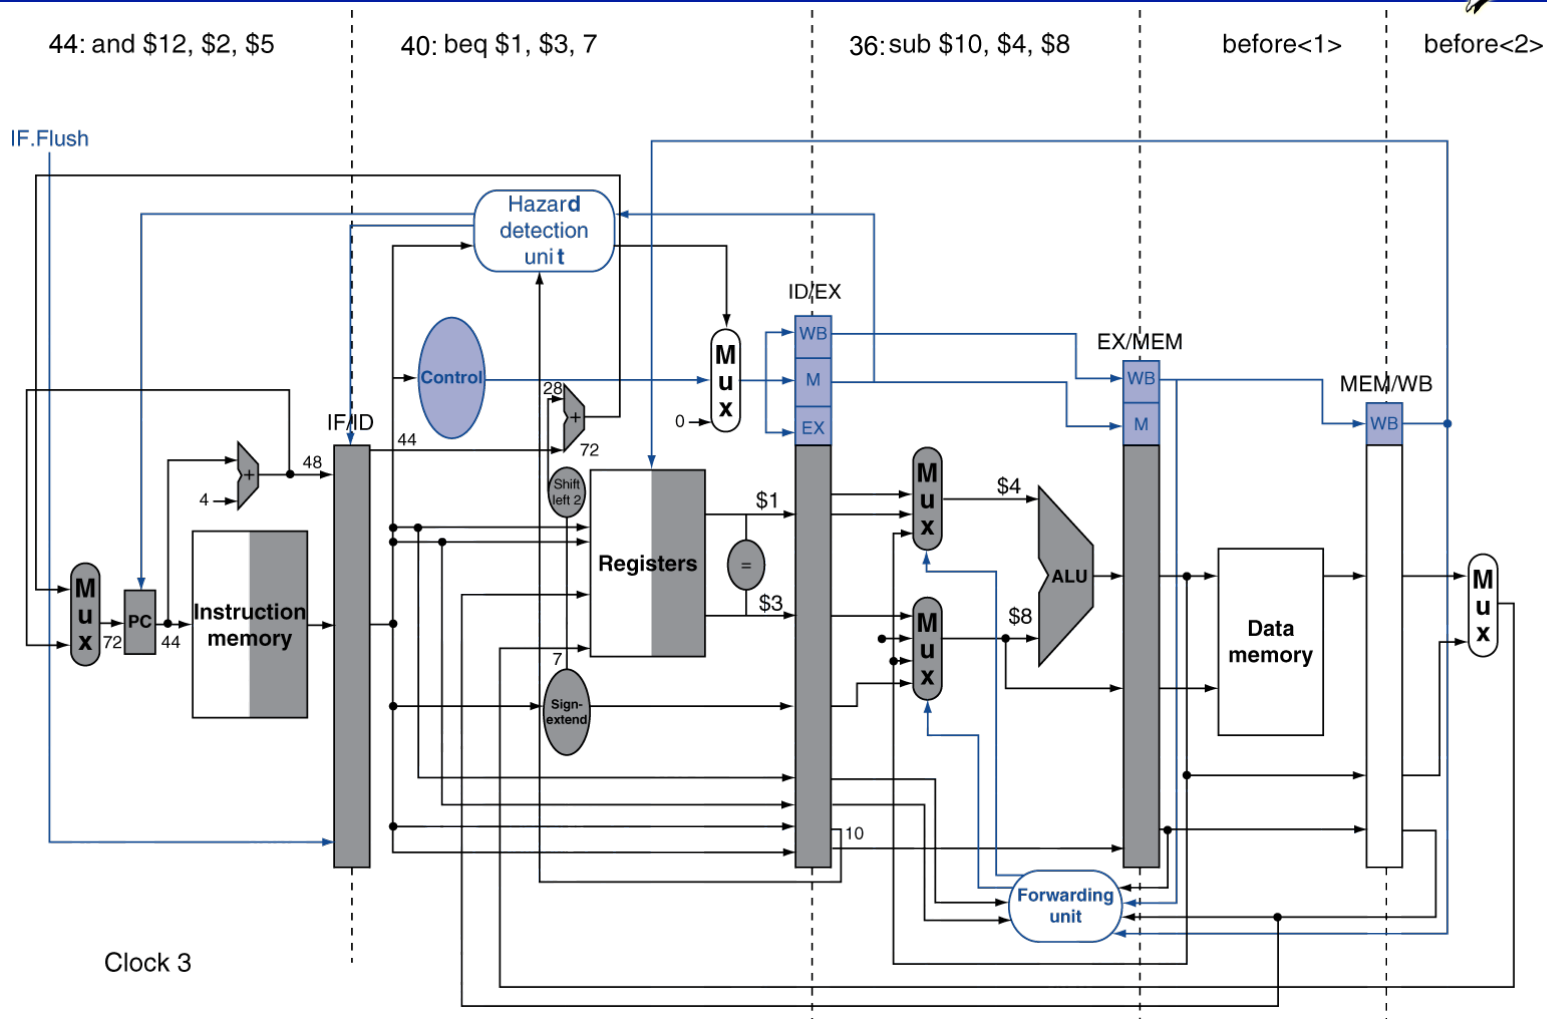
Reducing Branch Delay



move branch computation earlier in the pipeline
enable flush capability if mispredicted

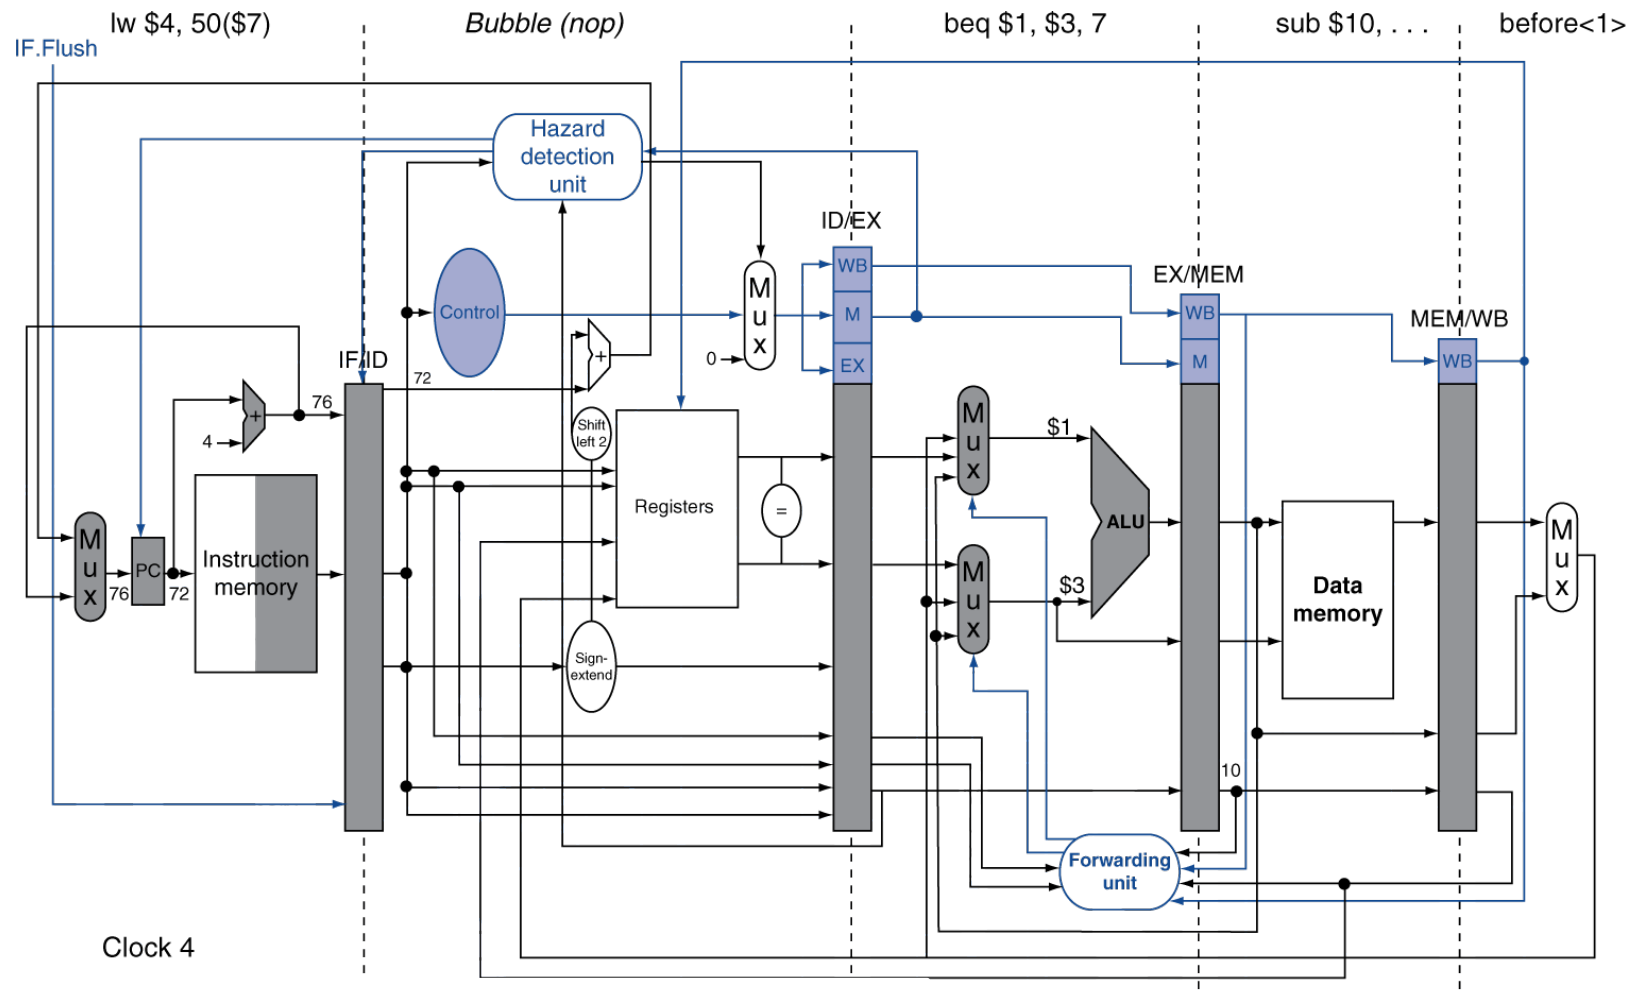


Example: Branch Taken





Example: Branch Taken



nop inserted using IF.Flush



How to Flush?

■ 5-stage Pipeline

- only one instruction to flush
- Can do by “flushing” the output register of IF

■ Deeper Pipelines

- Multiple cycles to determine branch outcome
- Need to know how many instructions in the pipe
- Flush all state changing actions
 - RegWrite, MemWrite, PCWrite (jmp, beq), overflow..

Performance Impact of Branch Stalls



- Need to stall for one cycle on every branch
- Consider the following case
 - The ideal CPI of the machine is 1
 - The branch causes a stall
- Effective CPI if 15% of the instructions are branches?
 - The new effective CPI is $1 + 1 \times 0.15 = 1.15$
 - The old effective CPI was $1 + 3 \times 0.15 = 1.45$



Delayed Branches

- An ISA-based solution used in early MIPS machines
 - Had a branch delay of one

- Example:

beq r1, r2, L

sub r4, r1, r3

and r6, r2, r7

Branch instruction

This operation ALWAYS is executed

This operation executes if branch fails

- Worked well initially
 - Compiler can fill one slot 50% of the time (50% nops)
- For modern processors, it is a pain
 - Many delay slots needed due to deeper pipelines
 - Processors deal with this using branch prediction
 - Delayed branches complicate exceptions as well




It's All Connected

- We moved the branch control point to ID
- Which means we consume two registers in ID
 - earlier in the pipeline as before (EX)
- What does this mean for forwarding/stalls?
 - Need to check!!



Data Hazards for Branches

- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction




- Can resolve using forwarding
 - Additional datapaths and control



Data Hazards for Branches


- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle





Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles





Branch CPI

15% Branch frequency

65% branches are taken

50% of single delay slots are filled usefully

Control point	Branch strategy	Branch Stall CPI
MEM	stall	
ID	stall	
ID	Predict taken	
ID	Predict Not taken	
ID (target) / MEM (==)	Predict taken	
ID	Delayed branch	



Branch Prediction

- Static prediction vs. dynamic prediction

- Static prediction schemes:

- always predict taken
- always predict not-taken
- compiler/programmer hint
- if (target < PC)
 - predict taken
- else
 - predict not-taken

What's the rationale behind this?



Dynamic Branch Prediction

- Branch History Table (BHT)
 - One entry for each branch PC
 - Taken/Not taken bit
- Branch Target Buffer (BTB)
 - One entry for each branch PC
 - Target address
- Increasingly important for long pipelines (IDx)
 - x86 vs. RISC-V instruction decode



Branch Prediction Example

```
void foo() {  
    for(e=0;e<4;e++) {  
        stuff ...  
    }  
}
```

```
foo:j E  
L:  stuff..  
   ..  
E:  bne $t3,$t4,L
```

How many times will bne at E: be predicted correctly?

Two-bit saturating counter:
strongly taken, weakly taken, weakly n-t, strongly n-t



Exceptions and Interrupts



Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
 - Switch from user to privileged (kernel) mode
 - Different ISAs use the terms differently
- Exception
 - Arises within the CPU
 - E.g., undefined opcode, overflow, syscall, ...
- Interrupt
 - E.g., from an external I/O controller (network card)
- Dealing with them without sacrificing performance is hard



Handling Exceptions

- Something bad happens to an instruction
 - Need to stop the machine
 - Fix up the problem (with privileged code like the OS)
 - Start the machine up again

- MIPS: managed with System Control Coprocessor (CP0)
 - Save PC of offending (or interrupted) instruction
 - In MIPS: Exception Program Counter (EPC)
 - Save indication of the problem
 - In MIPS: Cause register
 - We'll assume 1-bit for the examples
 - 0 for undefined opcode, 1 for overflow
 - Jump to handler code at hardwired address (e.g., 8000 00180)



An Alternate Mechanism

■ Vectored Interrupts

- Handler address determined by the cause
- Avoids the software overhead of selecting the code to run based on cause reg
- Jump table

■ Example:

- Undefined opcode: C000 0000
- Overflow: C000 0020
- ...: C000 0040

■ Instructions either

- Deal with the interrupt, or
- Jump to real handler



Handler Actions

- Read cause, and transfer to relevant handler
 - A software switch statement
 - May be necessary even if vectored interrupts are available
- Determine action required
- If restartable
 - Take corrective action
 - Use EPC to return to program
- Otherwise
 - Terminate program (e.g., segfault, ...)
 - Report error using EPC, cause, ...



Precise Exceptions

- Definition: precise exceptions
 - All previous instructions had completed
 - The faulting instruction was not started
 - None of the following instructions were started
 - No changes to the architecture state (registers, memory)
- Why are they desirable by OS developers?
- With single cycle design, precise exceptions are easy
 - Why?
- With pipelined design, they can be tricky
 - Why?



Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage
 - add \$1, \$2, \$1
 - Prevent \$1 from being clobbered by add
 - Complete previous instructions
 - Nullify subsequent instructions
 - Set Cause and EPC register values
 - Transfer control to handler
- Similar to mispredicted branch
 - Uses much of the same hardware

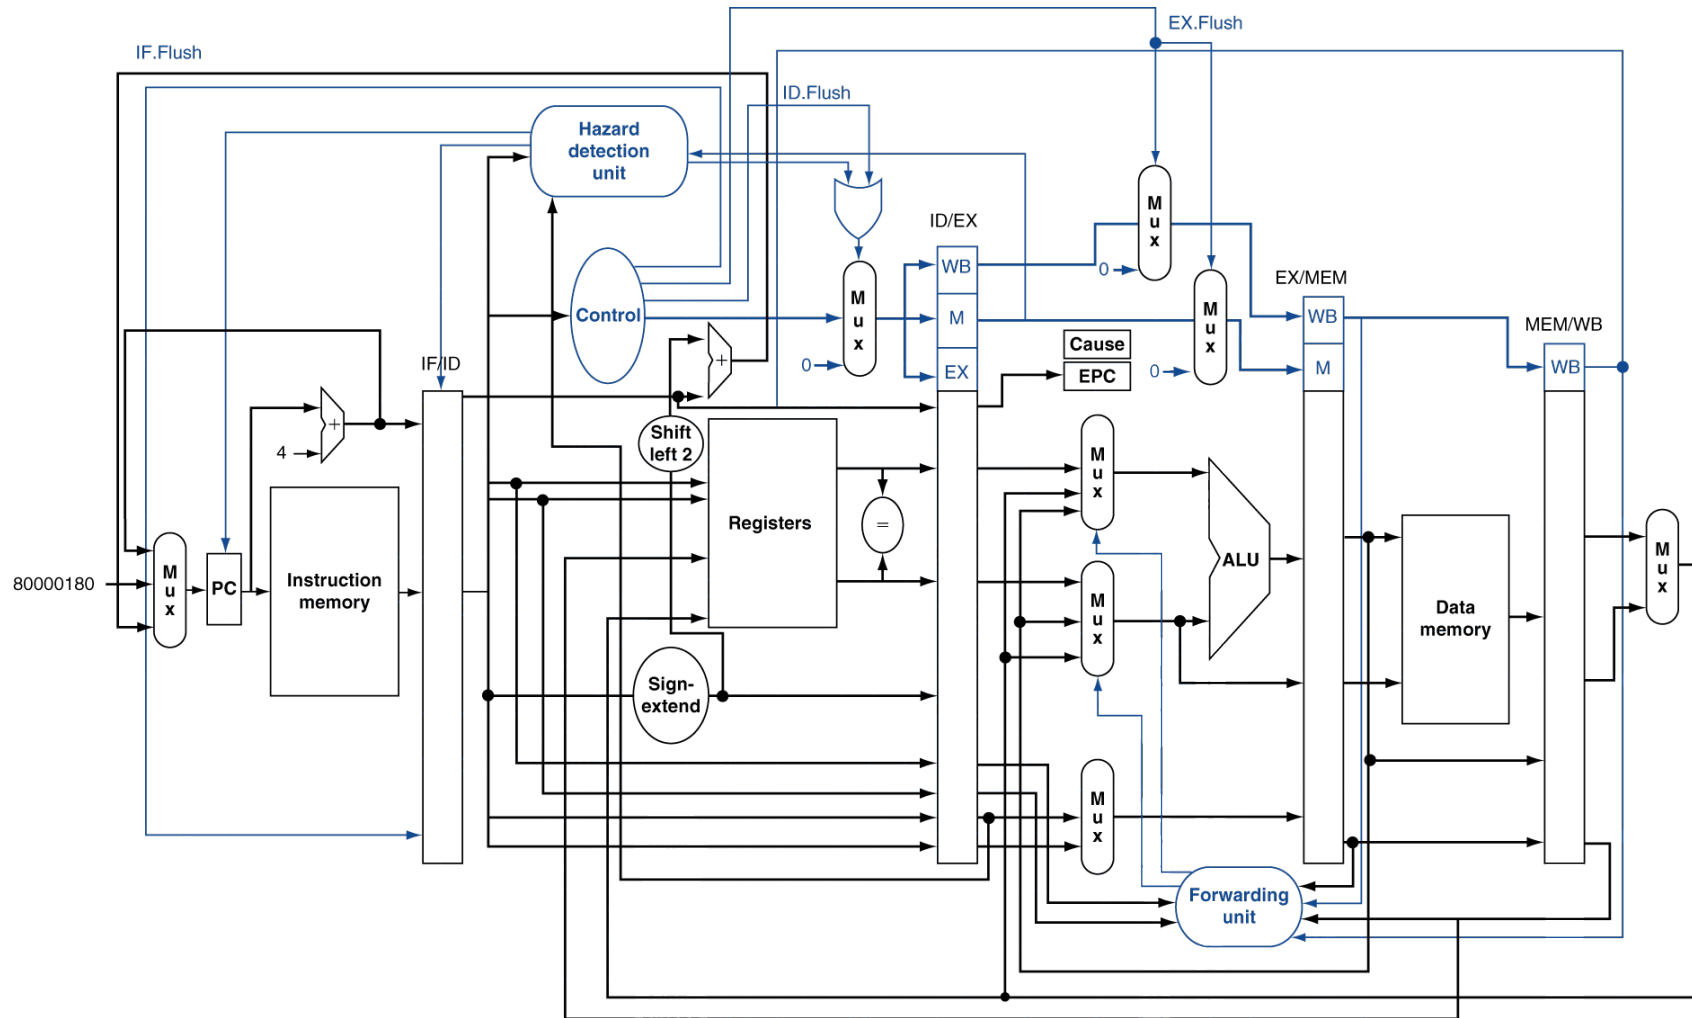


Nullifying Instructions

- Nullify = turn an instruction into a nop (or bubble)
 - Reset its RegWrite and MemWrite signals
 - Does not affect the state of the system
 - Resets its branch and jump signals
 - Does not cause unexpected flow control
 - Mark that it should not raise any exceptions of its own
 - Does not cause unexpected flow control
 - Let it flow down the pipeline



5-stage Pipeline with Exceptions





Exception Example

■ Exception on add in

40	sub	\$11,	\$2,	\$4
44	and	\$12,	\$2,	\$5
48	or	\$13,	\$2,	\$6
4C	add	\$1,	\$2,	\$1
50	s1t	\$15,	\$6,	\$7
54	lw	\$16,	50(\$7)	

...

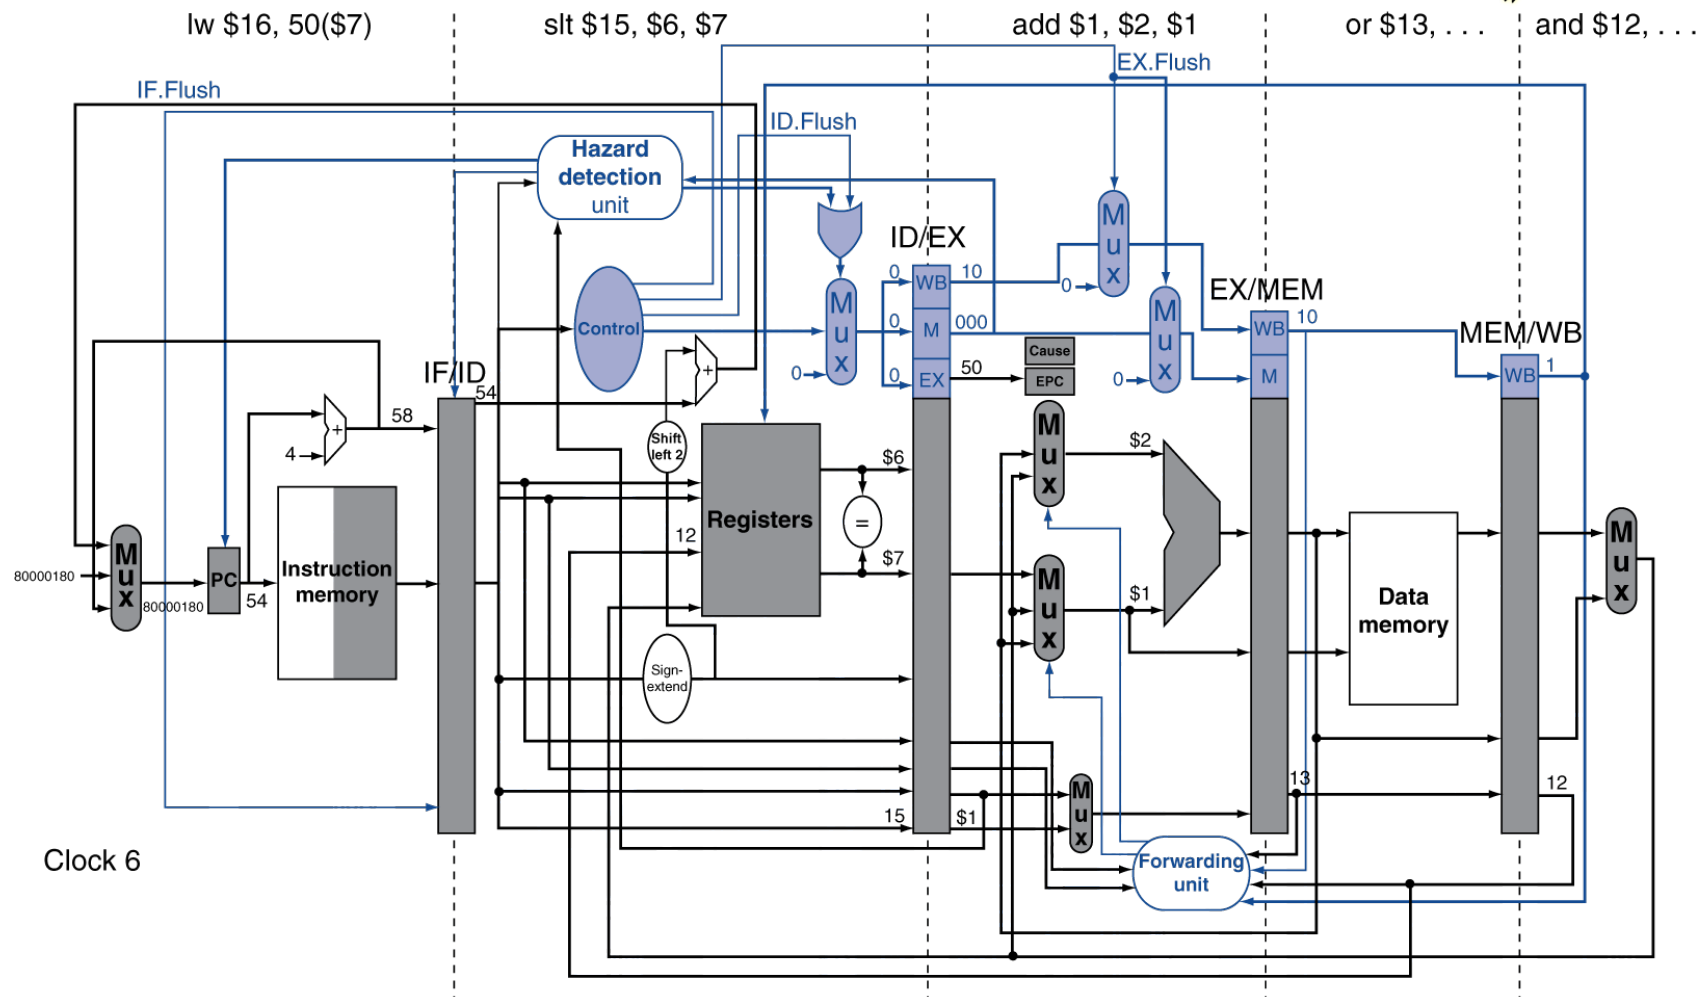
■ Handler

80000180	sw	\$25,	1000(\$0)
80000184	sw	\$26,	1004(\$0)

...

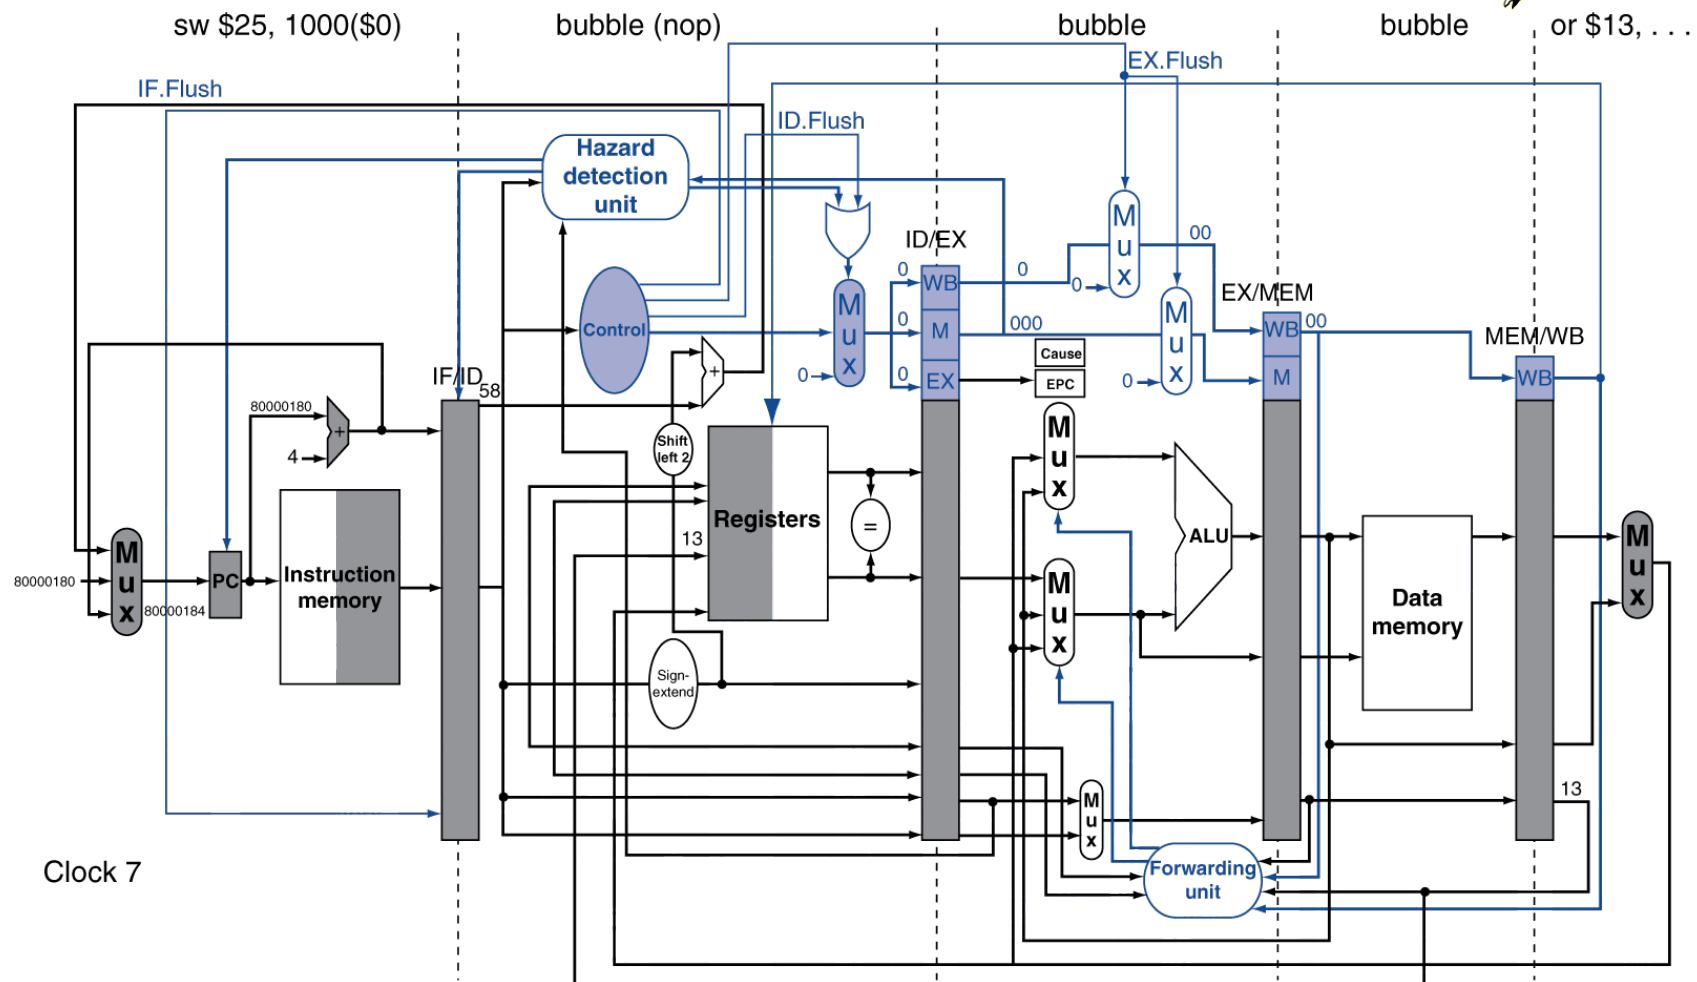


Exception Example





Exception Example





Dealing with Multiple Exceptions

- Pipelining overlaps multiple instructions
 - Could have multiple exceptions at once
- Simple approach: deal with exception in program order
 - Single commit point (MEM or WB stage)
 - Flush subsequent instructions
 - Necessary for “precise” exceptions



Beyond the 5-stage Pipeline

And why you should take CS316 😊

Advanced Multi-core Systems

Topics: advanced core design, GPU architecture, multi-threading, support for parallel programming, ...

And a research project



Beyond the 5-stage Pipeline

Desirable properties

Higher clock frequency

$CPI_{base} < 1$

Avoid in-order stalls

When instruction stalls, independent instructions behind it stall

Avoid stalls due to branches

How do we go about it?

Reminder: Key Tools for System Architects



1. Pipelining
2. Parallelism
3. Out-of-order execution
4. Prediction
5. Caching
6. Indirection
7. Amortization
8. Redundancy
9. Specialization
10. Focus on the common case



Modern Processors

High clock frequency → deep pipelines

10 – 20 stages

$CPI_{base} < 1$ → superscalar pipelines

Launch 2, 3, or 4 instructions every cycle

Avoid in-order stalls → out-of-order execution

Re-order instructions based on dependencies

Avoid stalls due to branches → branch prediction

Keep history about direction and target of branches



Why Does it Work: Instruction-level Parallelism

Independence among instructions

Example with ILP

add \$t0, \$t1, \$t2	Most programs have ILP
or \$t3, \$t1, \$t2	But it is not infinite
sub \$t4, \$t1, \$t2	Scope also matters a lot
and \$t5, \$t1, \$t2	

Example with no ILP

add \$t0, \$t1, \$t2
or \$t3, \$t0, \$t2
sub \$t4, \$t3, \$t2
and \$t5, \$t4, \$t2



Question

How much ILP is there in common programs?

Multiple Instruction Issue (Superscalar)



Fetch and execute multiple instructions per cycle \Rightarrow $CPI < 1$

Example: fetch 2 instructions/clock cycle

Dynamically decide if they can go down the pipe together or not (hazard)

Pipe Stages

Ideal CPI = 0.5	IF	ID	EX	MEM	WB		
	IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB	
		IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB
			IF	ID	EX	MEM	WB

Resources: double amount of hardware (FUs, Register file ports)

Issues: hazards, branch delay, load delay



Deeper Pipelining

Increase number of pipeline stages

Fewer levels of logic per pipe stage

Higher clock frequency

Example 9-stage pipeline

Pipe Stages

IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB	
	IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB

Issues: branch delay, load delay: $CPI = 1.4 - 2.0$, cost of pipeline registers

Modern pipelines

Pentium: 5 stages

Original Pentium Pro: 12 stages

Pentium 4: 31 stages, 3+ GHz

Tejas: 40-50 stages, 10 GHz, cancelled!

Core 2: 14 stages



Dynamic Scheduling

Execute instructions out-of-order

Fetch multiple instructions per cycle using branch prediction

Figure out which are independent and execute them in parallel

Example


```
add    $t0, $t1, $t2
or     $t3, $t0, $t2
sub    $t0, $t1, $t2
and    $t5, $t0, $t2
```

Superscalar + Dynamic scheduling

add \$t0, \$t1, \$t2	sub \$t0, \$t1, \$t2
or \$t3, \$t0, \$t2	and \$t5, \$t0, \$t2

What's wrong with this?

End of Uniprocessor Performance Improvements



Data collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, C. Batten