

# **CMPE110 Lecture 12**

## **Out-of-Order Execution**

---

Heiner Litz

<https://canvas.ucsc.edu/courses/12652>



# Announcements

- Midterm: May 7<sup>th</sup>
  - Everything covered in lectures, quizzes, HAs
  - Performance, ISA, processor architecture, 5-stage
  - Book Chapters: 2.1-2.8, 2.10, 2.14, 3.2, 4.1-4.10
  - Do I need to memorize the RISC-V ISA?
    - Midterm will have a copy of the cheat sheet
- Next Friday (May 4<sup>th</sup>):
  - No quiz
  - HA3 Review
  - Q & A

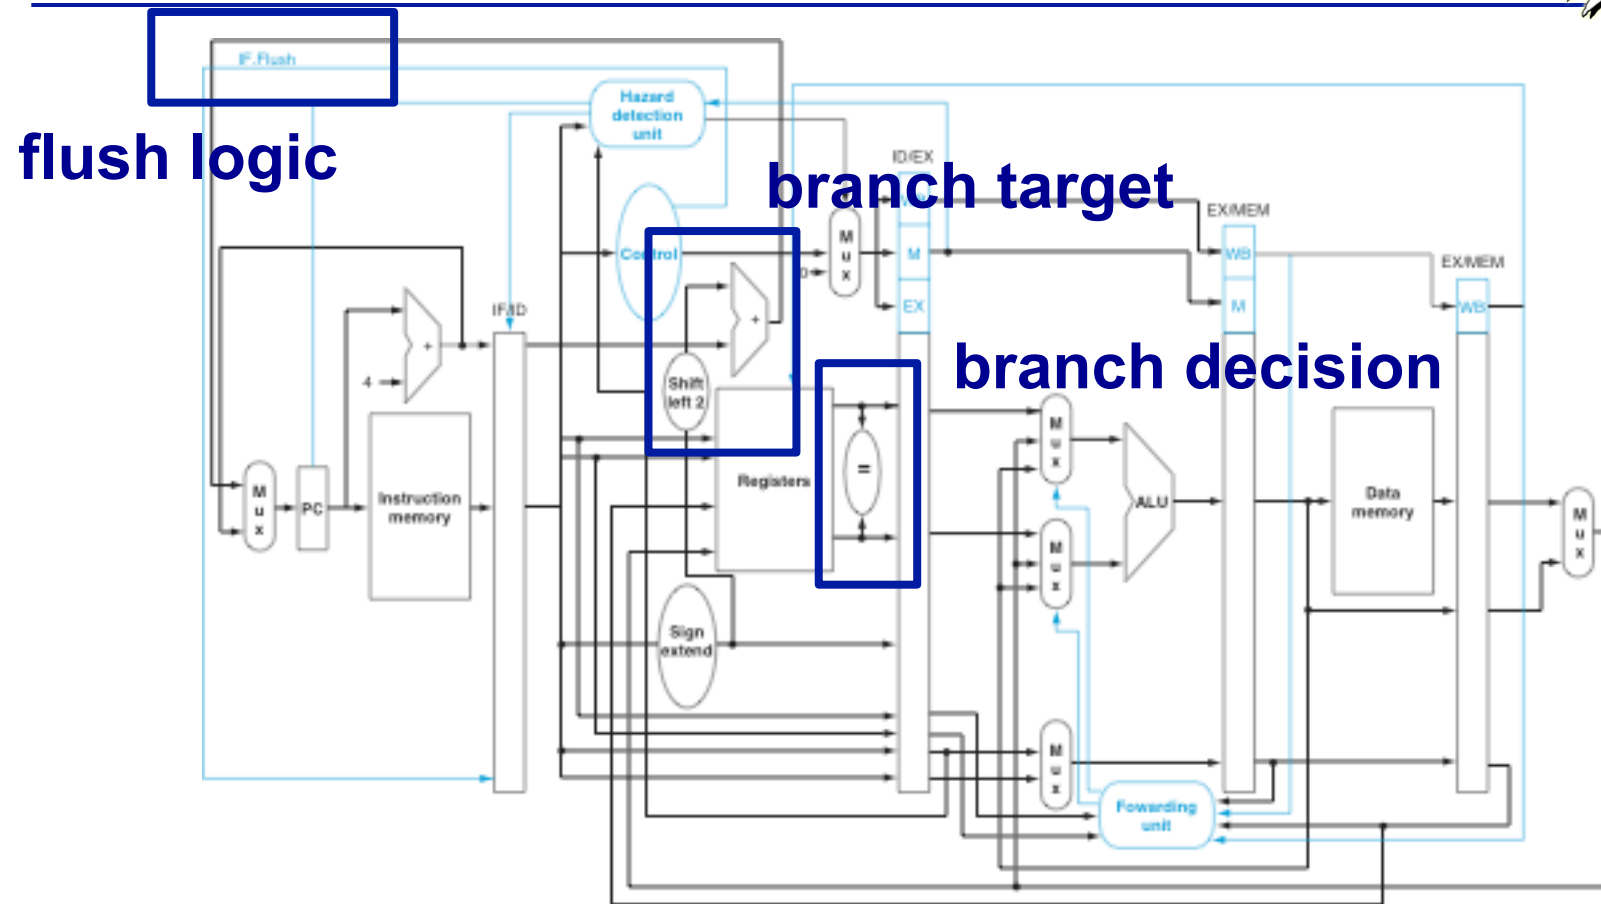
# Review

---





# Reducing Branch Delay




move branch computation earlier in the pipeline  
enable flush capability if mispredicted



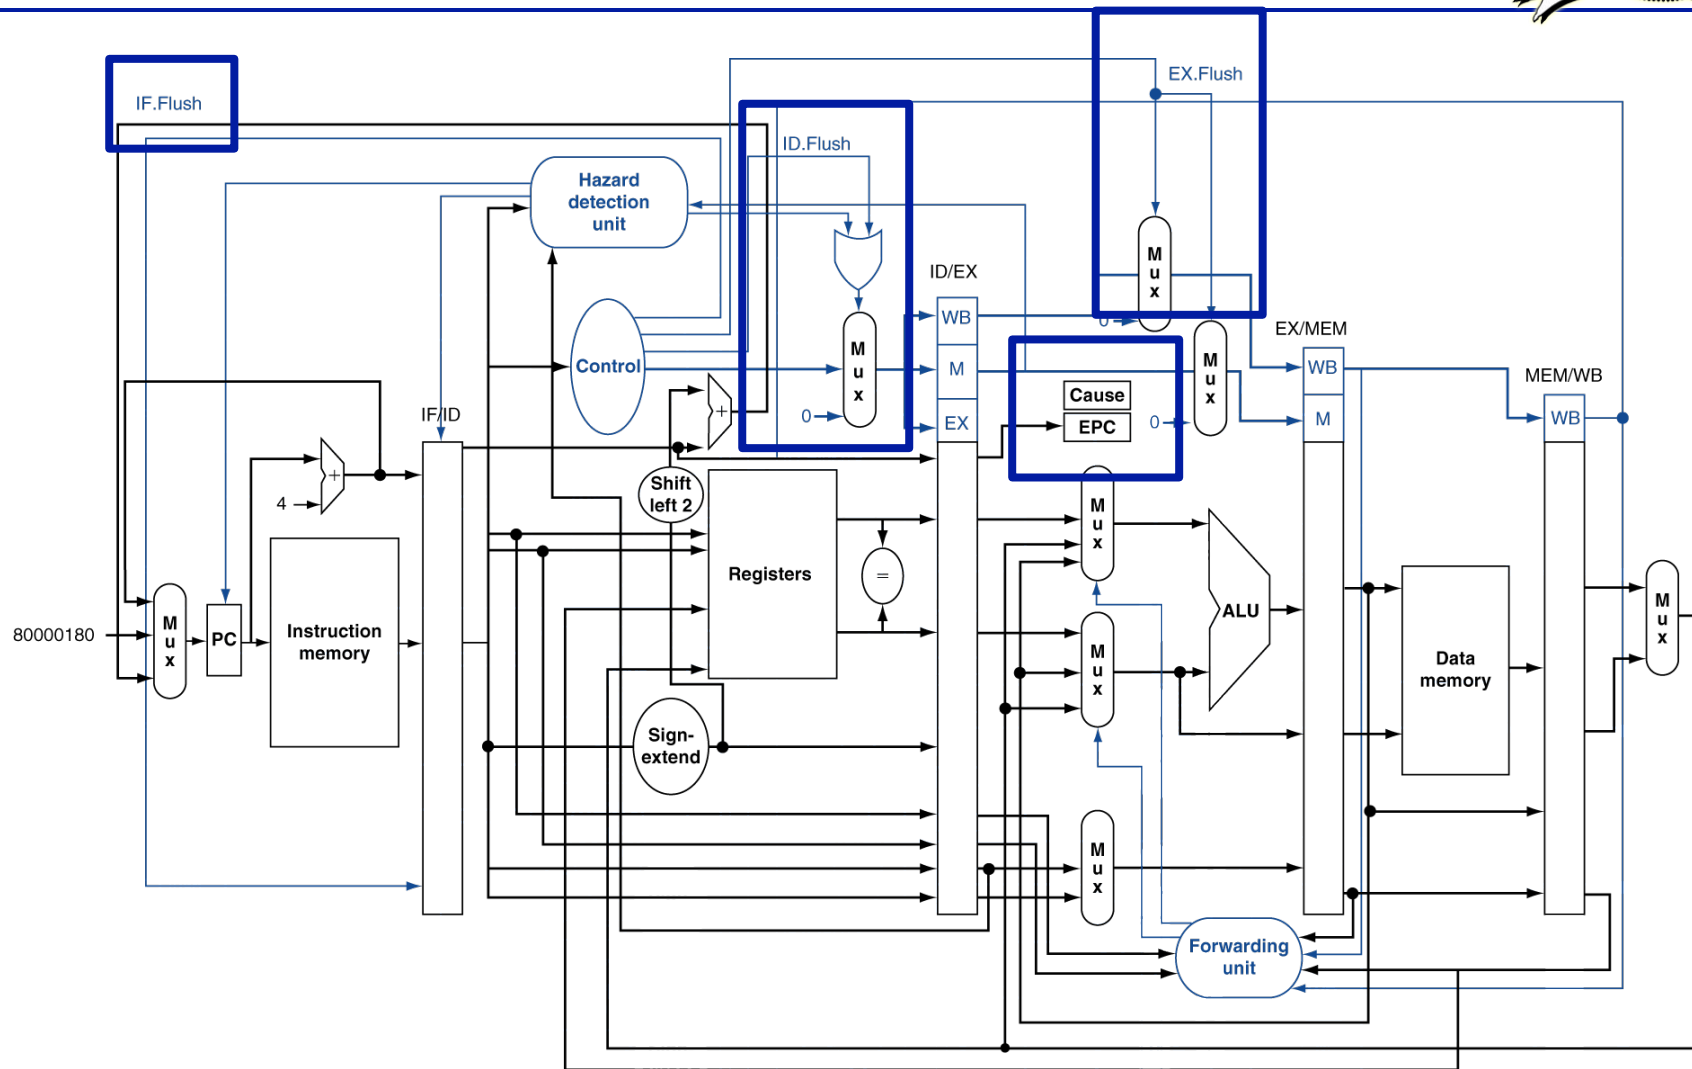
# Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
  - Need 2 stall cycles





# 5-stage Pipeline with Exceptions



# Beyond the 5-stage Pipeline

---





# Beyond the 5-stage Pipeline

---

You should take next Fall:

## Advanced Computer Architecture CMPE202 (Renau)

Topics: advanced core design, branch prediction, superscalar, symmetric multi-threading

## Advanced Parallel Processing CMPE220 (Litz) ☺

Topics: support for parallel programming, GPU architecture, SIMD, many core architectures, accelerators

And a research project





# Beyond the 5-stage Pipeline

## Desirable properties

Higher clock frequency

$CPI_{base} < 1$

Avoid in-order stalls

When instruction stalls, independent instructions behind it stall

Avoid stalls due to branches

## How do we go about it?



# Modern Processors

High clock frequency → deep pipelines

10 – 20 stages

$CPI_{base} < 1$  → superscalar pipelines

Launch 2, 3, or 4 instructions every cycle

Avoid in-order stalls → out-of-order execution

Re-order instructions based on dependencies

Avoid stalls due to branches → branch prediction

Keep history about direction and target of branches




# Instruction-Level Parallelism

$$D = 3(a - b) + 7ac$$

## ■ Data-flow execution order

## ■ Sequential execution order

ld a  
ld b  
sub a-b  
mul 3(a-b)  
ld c  
mul ac  
mul 7ac  
add 3(a-b)+7ac  
st d





# Question

---

How much ILP is there in common programs?

# Limits on Instruction Level Parallelism (ILP)



Weiss and Smith [1984]	1.58
Sohi and Vajapeyam [1987]	1.81
Tjaden and Flynn [1970]	1.86
Tjaden and Flynn [1973]	1.96
Uht [1986]	2.00
Smith et al. [1989]	2.00
Jouppi and Wall [1988]	2.40
Johnson [1991]	2.50
Acosta et al. [1986]	2.79
Wedig [1982]	3.00
Butler et al. [1991]	5.8
Melvin and Patt [1991]	6
Wall [1991]	7
Kuck et al. [1972]	8
Riseman and Foster [1972]	51
Nicolau and Fisher [1984]	90

- These studies made very different assumptions!
- Any guesses on the differences?

# Multiple Instruction Issue (Superscalar)



Fetch and execute multiple instructions per cycle  $\Rightarrow$   $CPI < 1$

Example: fetch 2 instructions/clock cycle

Dynamically decide if they can go down the pipe together or not (hazard)

## Pipe Stages

Ideal CPI = 0.5	IF	ID	EX	MEM	WB		
	IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB	
		IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB
			IF	ID	EX	MEM	WB

Resources: double amount of hardware (FUs, Register file ports)

Issues: hazards, branch delay, load delay

Modern processors: Intel x86: 4-way, ARM: 6-way/8-way



# Deeper Pipelining

Increase number of pipeline stages

Fewer levels of logic per pipe stage

Higher clock frequency

Example 9-stage pipeline

## Pipe Stages

IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB	
	IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB

Issues: branch delay, load delay:  $CPI = 1.4 - 2.0$ , cost of pipeline registers

Modern pipelines

Pentium: 5 stages

Original Pentium Pro: 12 stages

Pentium 4: 31 stages, 3+ GHz

Tejas: 40-50 stages, 10 GHz, cancelled!

Core 2: 14 stages



# Dynamic Scheduling

Execute instructions out-of-order

Fetch multiple instructions per cycle using branch prediction

Figure out which are independent and execute them in parallel

Example

```
add    $t0, $t1, $t2
or      $t3, $t0, $t2
sub     $t0, $t1, $t2
and     $t5, $t0, $t2
```

Superscalar + Dynamic scheduling

```
add    $t0, $t1, $t2    sub $t0, $t1, $t2
or      $t3, $t0, $t2    and $t5, $t0, $t2
```

What's wrong with this?





# Register Renaming

Rename (map) architectural registers to physical registers in decode stage to get rid of false dependencies

```
add    $t0, $t1, $t2
or      $t3, $t0, $t2
sub     $t0, $t1, $t2
and     $t5, $t0, $t2
```

Superscalar + Dynamic scheduling + register renaming

```
add    $t0A, $t1, $t2    sub $t0B, $t1, $t2
or      $t3, $t0A, $t2    and $t5, $t0B, $t2
```

Need more physical registers than architectural registers

Physical registers are automatically recycled

# How do we know which instructions can be scheduled?

---



- Need to represent the data flow graph in hardware

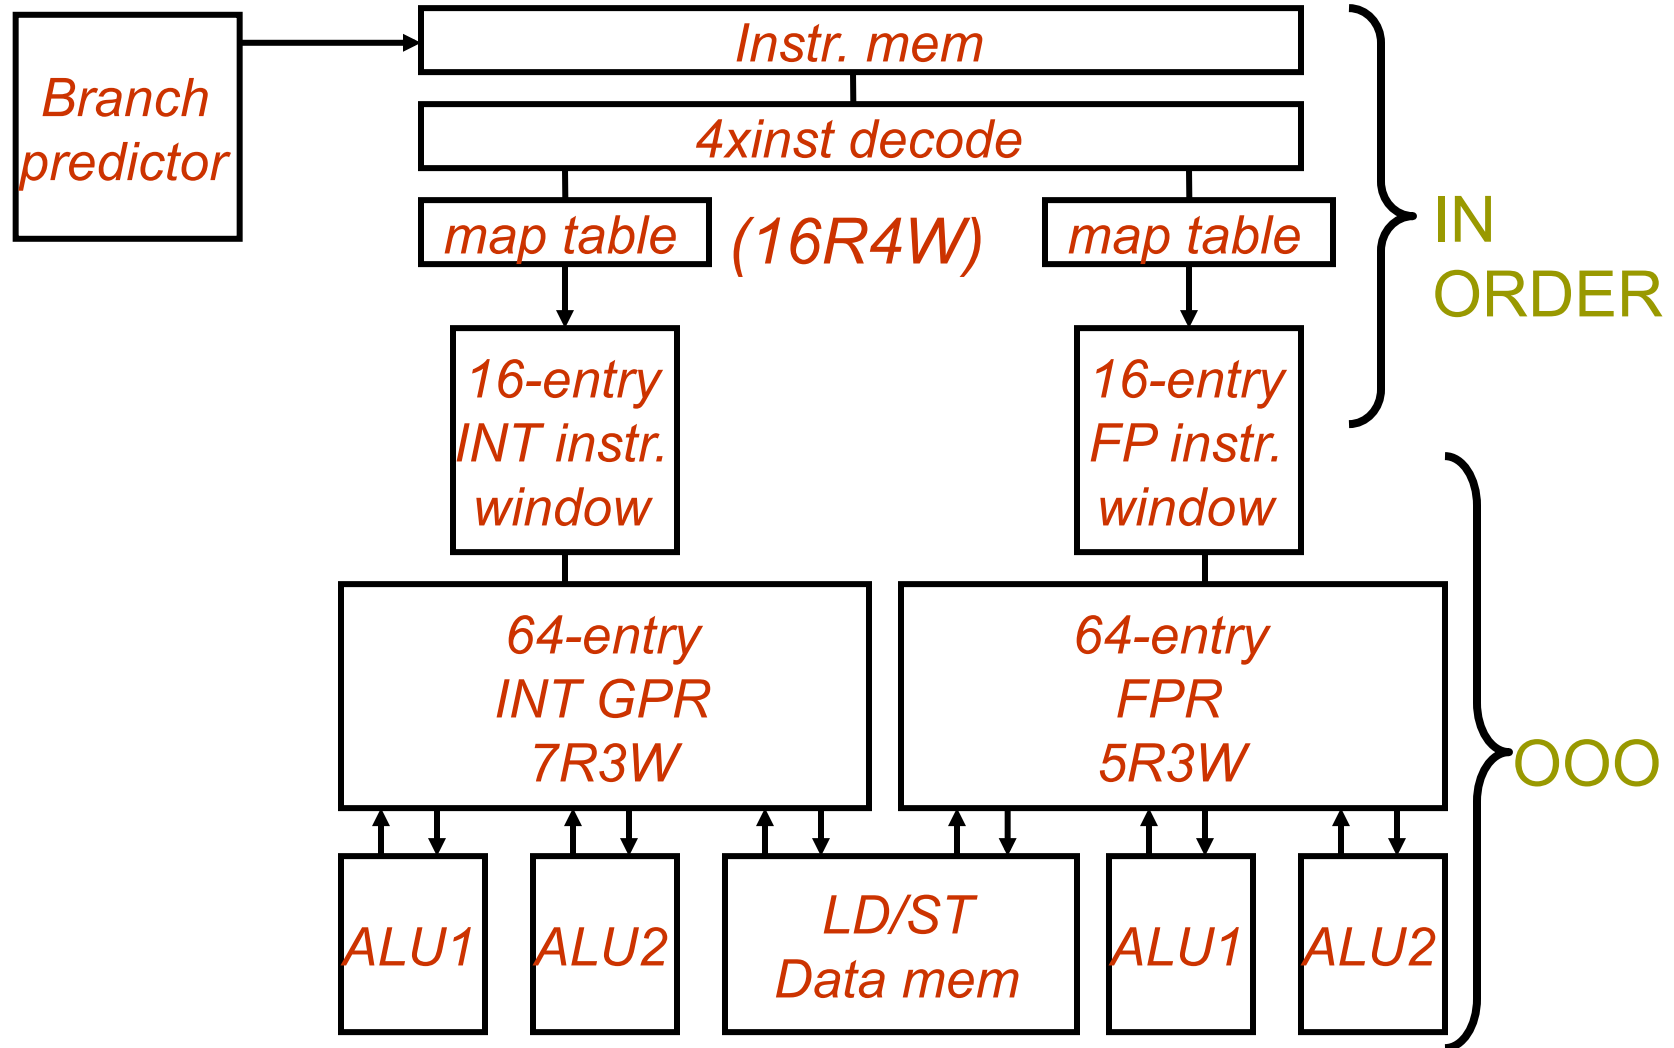


# Simple Scoreboarding

- Scoreboard: a bit-array, 1-bit for each GPR
  - if the bit is not set, the register has valid data
  - if the bit is set, the register has stale data
    - i.e. some outstanding inst is going to change it
- Dispatch in order:  $RD \leftarrow F_n(RS, RT)$ 
  - if  $SB[RS]$  or  $SB[RT]$  is set  $\Rightarrow$  RAW, stall
  - if  $SB[RD]$  is set  $\Rightarrow$  WAW, stall
  - else dispatch to FU, set  $SB[RD]$ 
    - Assuming a functional unit is available (structural hazard)
- Complete out-of-order
  - update  $GPR[RD]$ , clear  $SB[RD]$

Can be avoided  
with renaming



# Dynamic Scheduling in a Modern OOO MIPS R10000





# R2000 vs. R10000

- R2000: mid-80's, 32b 5-stage integer pipeline only, 120 K transistors
- R10k: mid-90's, 64b out-of-order superscalar w/ integer & FPU, 64 KB cache, 6.8 M transistors







# Limits of Advanced Pipelining

Limited ILP in real programs

Pipeline overhead

Branch and load delays exacerbated

Clock cycle timing limits

Limited branch prediction accuracy (85%-98%)

Even a few percent really hurts with long/wide pipes!

Memory inefficiency


Load delays + # of loads/cycle

Caches (next lecture) are not perfect

Complexity of implementation

Expensive to design

Why does a superscalar, out-of-order processor burn more power and energy?



# Caveats of Out-of-Order Processors

---



What can go wrong if I reorder instructions?

## Tips

Think of dependencies

Think of exceptions

Solutions?



# Performance/Power Tradeoffs


Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Yes	Yes
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level, simple	2-level, aggressive
1 <sup>st</sup> level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 <sup>nd</sup> level caches/core	128-1024 KiB	256 KiB
3 <sup>rd</sup> level caches (shared)	-	2- 8 MB






# Performance/Power Tradeoffs

## ARM A8




Minnespec Benchmarks

## Intel i7 920



SPEC 2006 integer benchmarks


# Impact of Branch Missprediction



**FIGURE 4.79** Percentage of branch mispredictions and wasted work due to unfruitful speculation of Intel Core i7 920 running SPEC2006 integer benchmarks.



# End of Uniprocessor Performance Improvements



Data collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, C. Batten