

# **CMPE110 Lecture 24**

## **Multicore**

---

Heiner Litz

<https://canvas.ucsc.edu/courses/12652>

# Announcements



## ■ Next week

- Monday: Pat will give exam review session 1
- Wednesday: Heiner will give Zoom lecture on Multicore/GPU
- Friday: Heiner will give exam review session 2

## ■ Office Hours

- No office hours on Mon 4<sup>th</sup> and Wed 6<sup>th</sup>
- Office hours today 2:30-4pm
- Office hours Friday 8<sup>th</sup>, 2:30-4pm

## Additional Reading Material

---




- I/O: Computer Architecture, A Quantitative Approach, 4<sup>th</sup> Edition, Chapter 6, Storage Systems
- Multicore: Computer Organization and Design, Chapter 6, Parallel Processors, 6.1 – 6.5

# Review

---

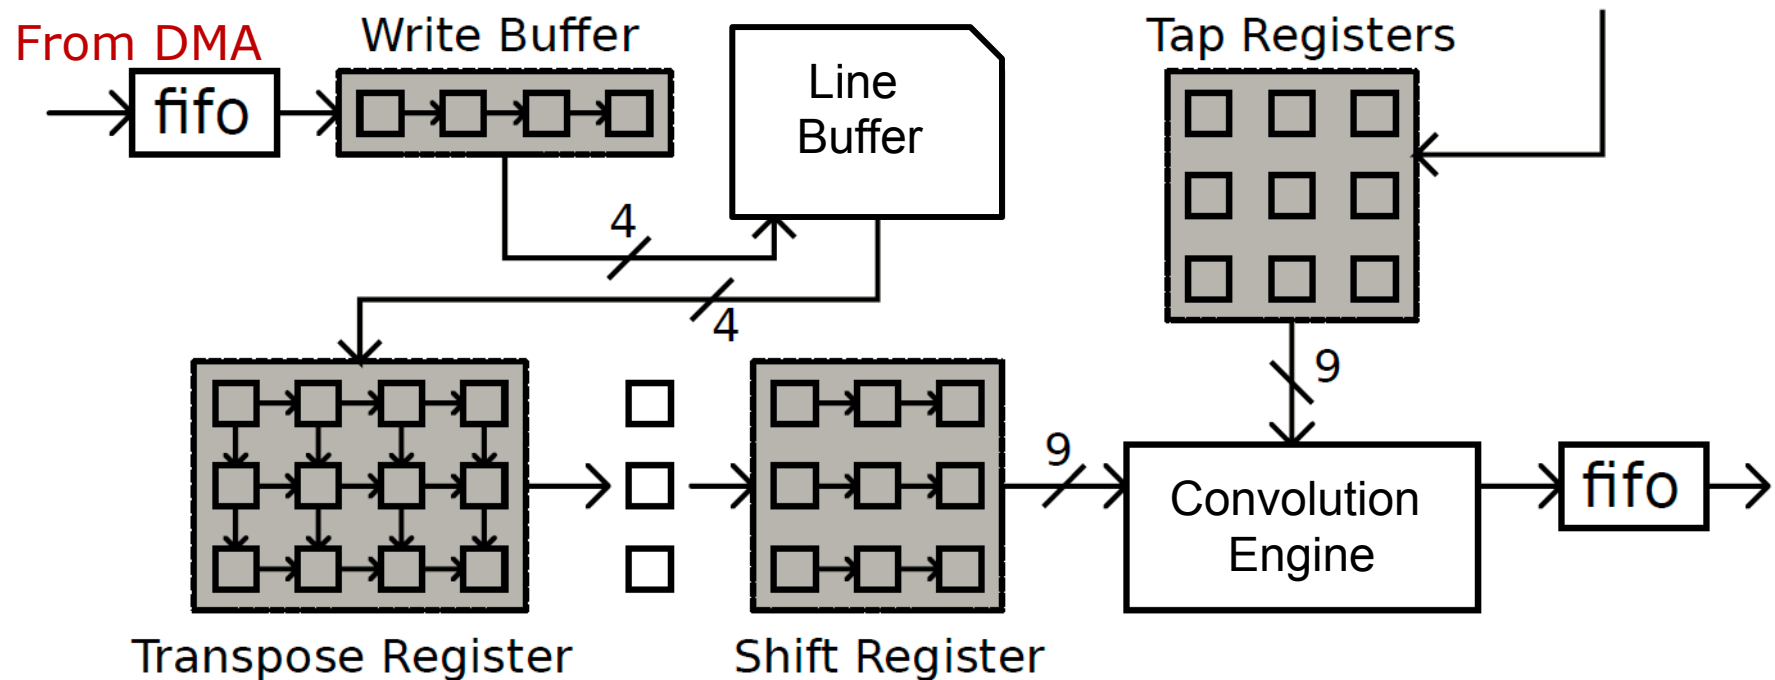


# Convolution Engine




- How many rows of data do we need to buffer?
- How much new data do we need per pixel?

# Convolution Datapath

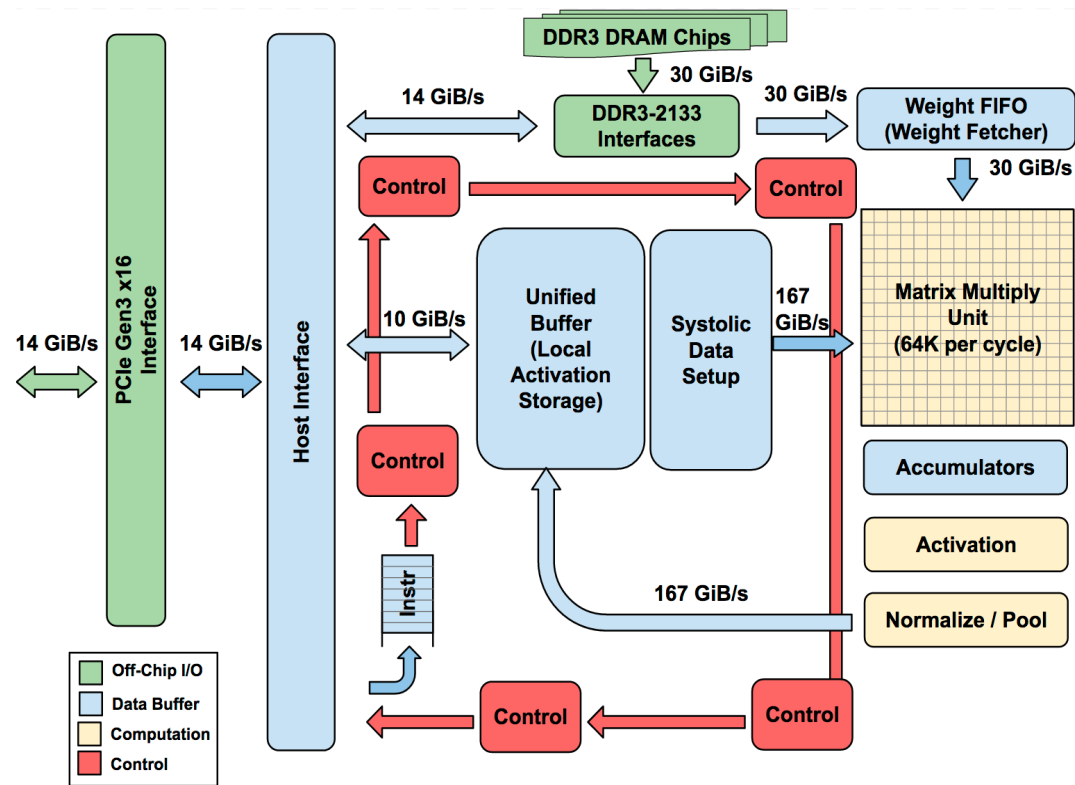


- How many pixels per cycle?

# Example - FPGAs



# Google TPU






# Multicore

---



# Sequential Processor Performance has Plateaued




Data collected by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, C. Batten



## Why does Multicore help?

- Increasing core count increases power, right?
- Pollack's rule: Performance increase due to microarchitecture improvements equals  $\sqrt{\text{complexity}}$



# Parallel Computing



Goal: higher performance from multiple processors

Multiprocessors → multi-core

Scalability, availability, power efficiency

## Challenge

Making effective use of parallel hardware

## Process-level parallelism (Multi-process)

High throughput for independent jobs

## Parallel program (Multi-threading)

Single program run on multiple processors



# What We've Already Covered

---

## Parallelism and ISA

- Instruction Level Parallelism (ILP)

- Vector parallelism (SIMD or data-level)

## Parallelism and memory hierarchy

- Cache coherence

# Parallel Programming



Parallel software is the problem

Need to get significant performance improvement

Otherwise, just use a faster uniprocessor (easier)

Unless you can't find a faster uniprocessor!

## Difficulties

Partitioning

Coordination

Communications overhead

# Amdahl's Law



Sequential part can limit speedup

Example: 100 processors, 90× speedup?

$$T_{\text{new}} = T_{\text{parallelizable}}/100 + T_{\text{sequential}}$$

$$\text{Speedup} = \frac{1}{(1 - F_{\text{parallelizable}}) + F_{\text{parallelizable}}/100} = 90$$

$$\text{Solving: } F_{\text{parallelizable}} = 0.999$$

Need sequential part to be 0.1% of original time



## Scaling Example

Workload: sum of 10 scalars, and  $10 \times 10$  matrix sum

Speed up from 10 to 100 processors

Single processor: Time =  $(10 + 100) \times t_{\text{add}}$

10 processors

$$\text{Time} = 10 \times t_{\text{add}} + 100/10 \times t_{\text{add}} = 20 \times t_{\text{add}}$$

$$\text{Speedup} = 110/20 = 5.5 \text{ (55\% of potential)}$$

100 processors

$$\text{Time} = 10 \times t_{\text{add}} + 100/100 \times t_{\text{add}} = 11 \times t_{\text{add}}$$

$$\text{Speedup} = 110/11 = 10 \text{ (10\% of potential)}$$

Assumes load can be balanced across processors





## Scaling Example (cont)

What if matrix size is  $100 \times 100$ ?

Single processor: Time =  $(10 + 10000) \times t_{\text{add}}$

10 processors

$$\text{Time} = 10 \times t_{\text{add}} + 10000/10 \times t_{\text{add}} = 1010 \times t_{\text{add}}$$

$$\text{Speedup} = 10010/1010 = 9.9 \text{ (99\% of potential)}$$

100 processors

$$\text{Time} = 10 \times t_{\text{add}} + 10000/100 \times t_{\text{add}} = 110 \times t_{\text{add}}$$

$$\text{Speedup} = 10010/110 = 91 \text{ (91\% of potential)}$$

Assuming load balanced



## Strong vs Weak Scaling

Strong scaling: problem size fixed

As in first example

Weak scaling: problem size proportional to number of processors

10 processors,  $10 \times 10$  matrix

$$\text{Time} = 20 \times t_{\text{add}}$$

100 processors,  $32 \times 32$  matrix

$$\text{Time} = 10 \times t_{\text{add}} + 1000/100 \times t_{\text{add}} = 20 \times t_{\text{add}}$$

Constant performance in this example

# Flynn's Taxonomy: Classification of Parallel Machines



		Data Streams	
		Single	Multiple
Instruction Streams	Single	<b>SISD:</b> Intel x86 processors	<b>SIMD:</b> SSE instructions of x86
	Multiple	<b>MISD:</b> No examples today	<b>MIMD:</b> Any multicore


## SPMD: Single Program Multiple Data

A parallel program on a MIMD computer

Conditional code for different processors based on data



# Shared Memory Multi-core



- Shared memory multi-core
  - Each core has private PC and registers
  - All cores access the same, shared memory
  - Communication through loads and stores
  - Synchronization through atomic instructions
- Requires cache coherence



## Example: Sum Reduction

- Sum 100,000 numbers on 100 processors

- Each processor has ID:  $0 \leq P_n \leq 99$
- Partition 1000 numbers per processor
- Initial summation on each processor

```
sum[Pn] = 0;
```

```
for (i = 1000*Pn; i < 1000*(Pn+1); i = i + 1)  
    sum[Pn] = sum[Pn] + A[i];
```


- Now need to add these partial sums

- Reduction: divide and conquer
- Half the processors add pairs, then quarter, ...
- Need to synchronize between reduction steps

# Example: Sum Reduction



```
half = 100;
repeat
  synch();
  if (half%2 != 0 && Pn == 0)
    sum[0] = sum[0] + sum[half-1];
    /* Conditional sum needed when half is odd;
       Processor0 gets missing element */
  half = half/2; /* dividing line on who sums */
  if (Pn < half) sum[Pn] = sum[Pn] + sum[Pn+half];
until (half == 1);
```



# Synchronization



- Assume multiple threads on multiple cores
- Shared data structure, e.g. a linked-list
- What happens when multiple threads work on the same list?
- How do we solve this issue?



## Critical Sections (Locking)

```
bool void delete(item) {  
    lock_list();  
    found = lookup_item(item);  
    if (found)  
        delete_item(item);  
    unlock_list();  
}
```

**What is the problem with this approach?**



## Solution? Fine grain locking



```
bool void delete(item) {  
    found = lookup_item(item);  
    if (found)  
        lock_prev_item();  
        delete_item(item);  
        unlock_prev_item();  
}
```



# What do we need to implement locks?

```
bool lock = 0;
```

thread a:

```
if(lock==0)
```

```
lock = 1
```

```
//enter critical section
```

time



thread b:

```
if(lock==0)
```

```
lock = 1
```

```
//enter critical section
```

Atomic read-modify-write (RMW) instructions



## Atomic Instructions in RISC-V

- load reserved (lr.w)
- store conditional (st.w)
- Example: (lock variable's memory address stored in x3)

retry:

lr.w x2, 0(x3)	//load lock variable into reg x2
beq, x2, 1, retry	//check if already locked
addi x2, x2, 1	//set lock variable in reg
st.w x5, x2, 0(x3)	//store lock var, success: x5 <- 0
bne x5, x0, retry	//check whether store was successful

# Multithreading




- Idea: execute multiple threads on a single core (or threads > cores)
- Threads within a process
  - Separate: PC, registers
  - Shared: virtual memory space
- Key advantage
  - When 1 thread stalls, switch to another thread
  - Stalls: pipeline, memory, I/O, ...
  - The faster the switch, the more stall events we can mask

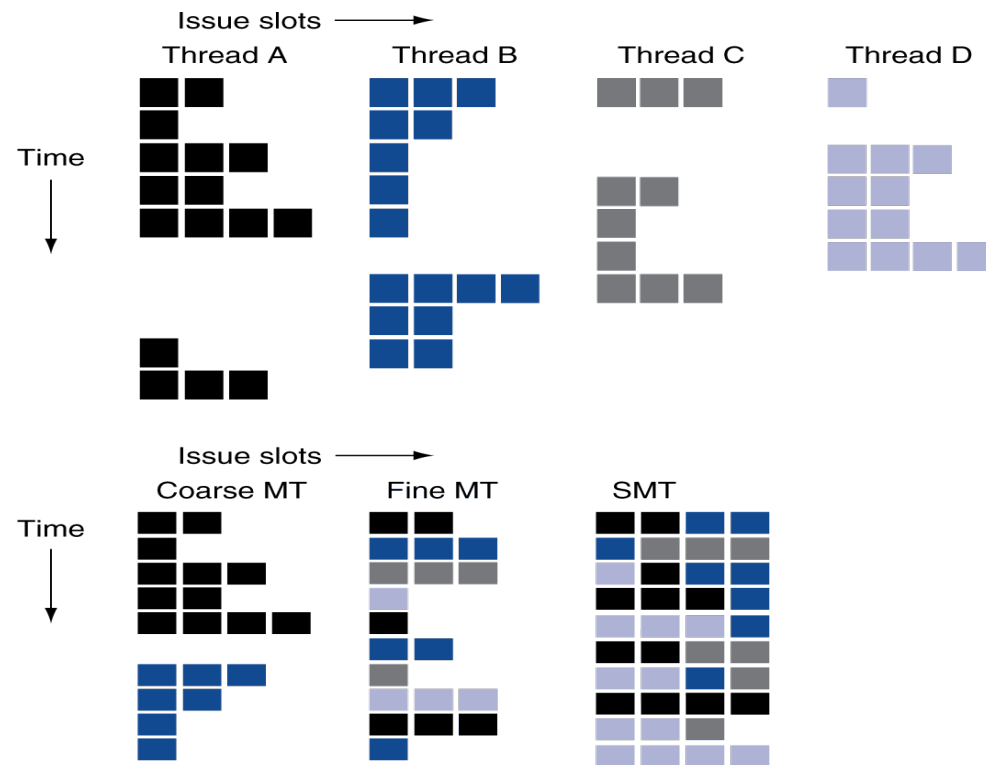
## Reminder: Superscalar Processor



- Superscalar: Process  $>1$  instruction per cycle
- Effectively multiple pipelines in parallel
- Exploits ILP by leveraging control flow graph
- Perform dynamic dependency checking



# Types of Multithreading



# Types of Multithreading



## ■ Fine-grain multithreading

- Switch threads after each cycle
- Interleave instruction execution
- If one thread stalls, others are executed

## ■ Coarse-grain multithreading

- Only switch on long stall (e.g., L2-cache miss)
- Simplifies hardware, but doesn't hide short stalls (e.g., data hazards)

# Simultaneous Multithreading



- In multiple-issue dynamically scheduled processor
  - Schedule instructions from multiple threads
  - Instructions from independent threads execute when function units are available
  - Within threads, dependencies handled by scheduling and register renaming
- Example: Intel Pentium-4 HT
  - Two threads: duplicated registers, shared function units and caches



# Concluding Remarks



- Goal: higher performance by using multiple processors
- Difficulties
  - Developing parallel software
  - Devising appropriate architectures
- Options
  - Vectors/SIMD, multi-core, GPUs, clusters
- Important notes
  - SaaS importance is growing and clusters are a good match
  - Performance per \$ and Joule drive both mobile and WSC